

# BIDS-Training 2024

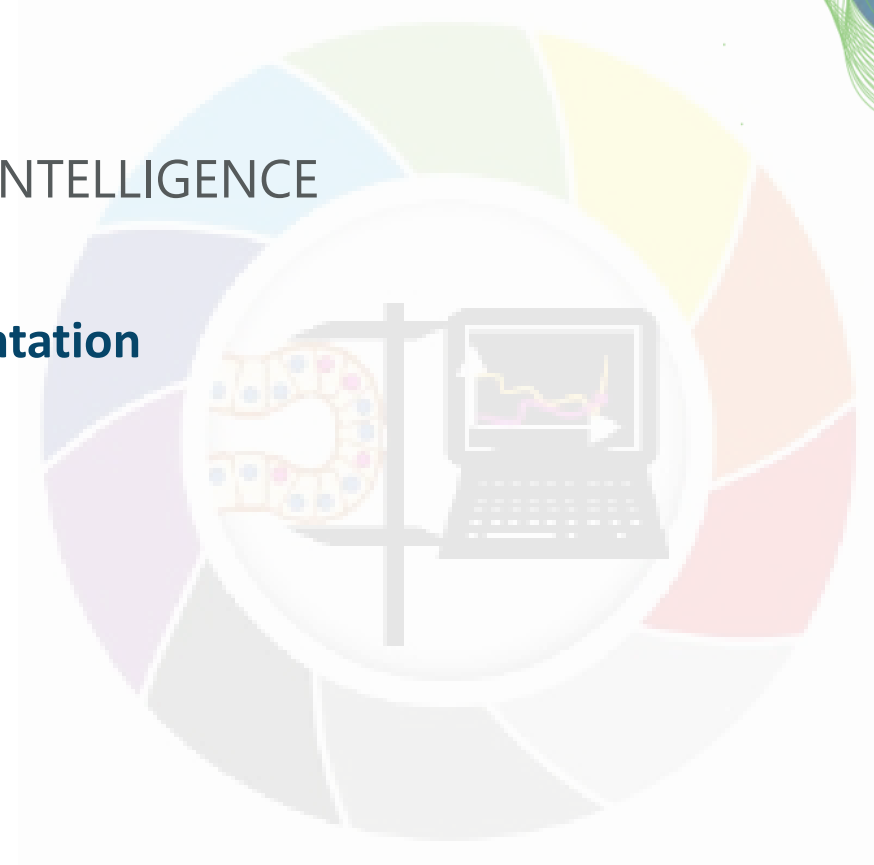


CENTER FOR SCALABLE DATA ANALYTICS AND ARTIFICIAL INTELLIGENCE

**Day 2, Session 3: Machine Learning for Pixel and Object Segmentation**

**SPEAKER: Christian Martin, Anja Neumann**

**DATE: 14-05-2024**



These slides can be reused under the [CC-BY 4.0](https://creativecommons.org/licenses/by/4.0/) license unless mentioned otherwise.

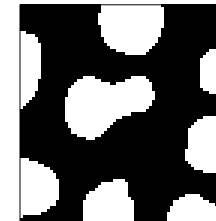
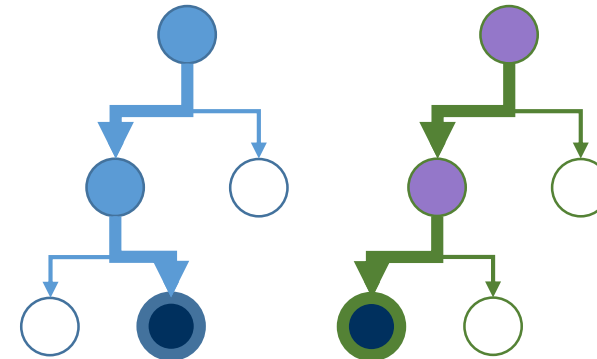
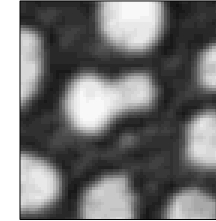
# Overview

## Machine Learning

- Introduction
- Decision Tree / Random Forest
- Image Segmentation using thresholding
- Image Segmentation using machine learning
- Object classification
- Segmentation quality
- Model validation
- Outlook

## Practical part with Python

- scikit-learn / napari
- Accelerated pixel and object classification (APOC)

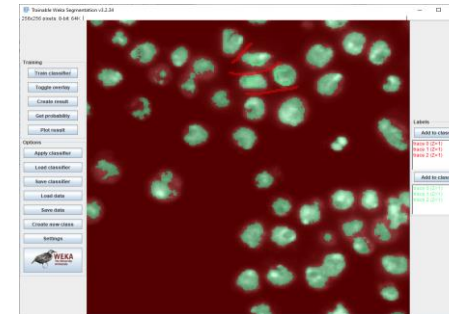


# Machine learning

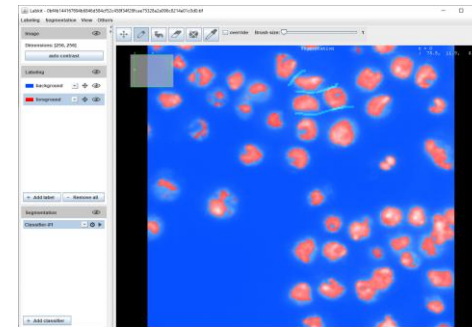
- A research field in computer science
- Finds more and more applications, also in life sciences.

Artificial intelligence

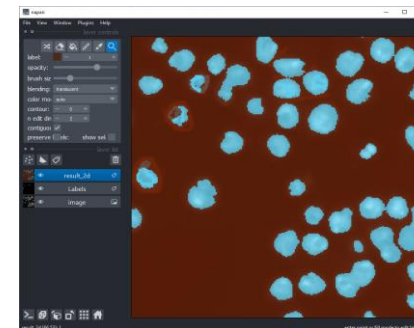
Machine learning



Trainable Weka Segmentation  
<https://imagej.net/plugins/tws/>



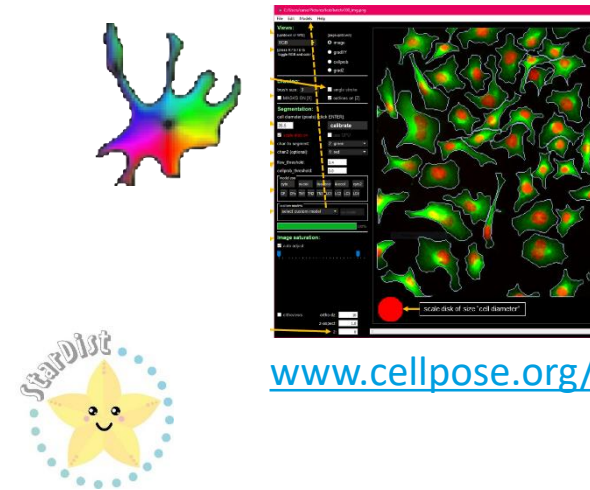
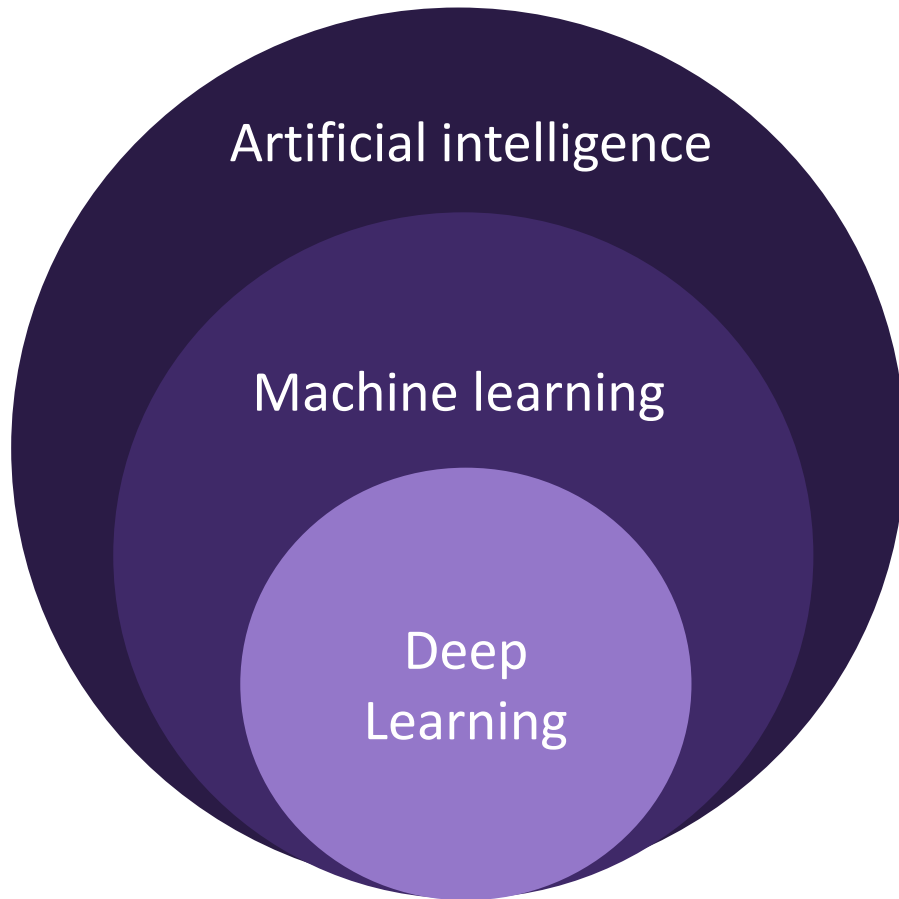
LabKit  
<https://imagej.net/plugins/labkit/>



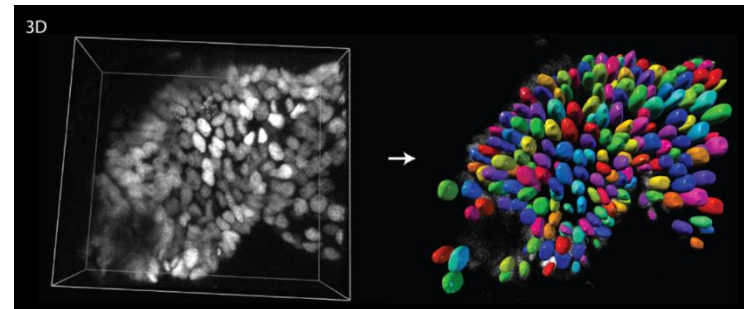
Python/scikit-learn/napari/apoc

# Machine learning

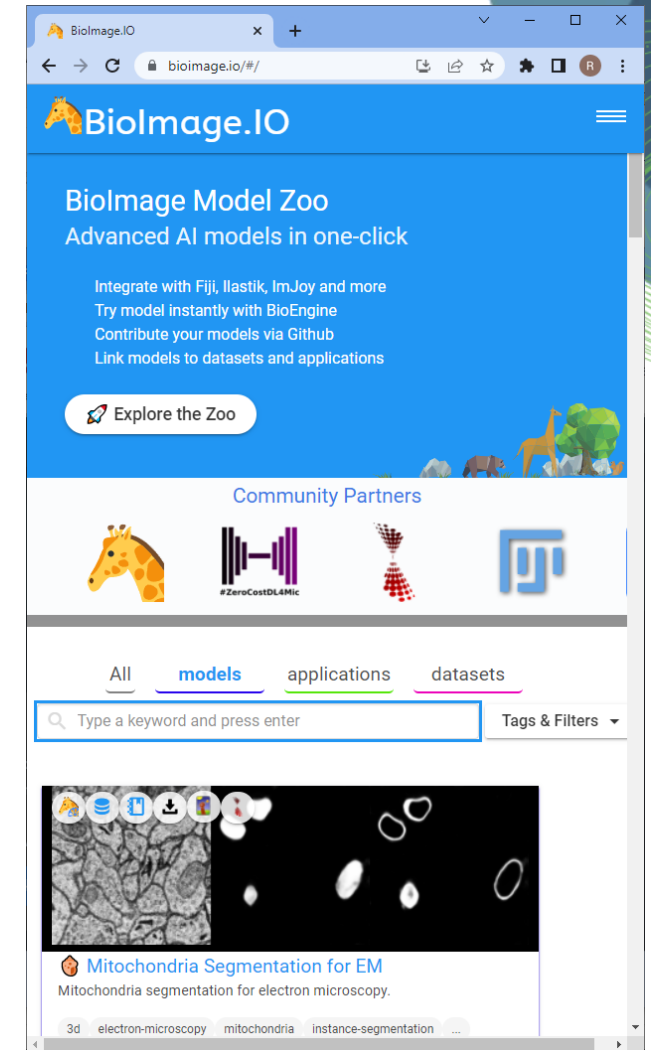
- A research field in computer science
- Finds more and more applications, also in life sciences.



[www.cellpose.org/](https://www.cellpose.org/)



<https://github.com/stardist/stardist>



<https://bioimage.io/>

Logos and screenshots are taken from the github repositories / websites provided under BSD and MIT licenses.

# Machine Learning

## Machine Learning

- subfield of Artificial Intelligence
- Automatic construction of predictive models from given data
- Learning from Data (data-driven approach)
- Input Data: m items of n dimensions
- If available, ground truth for each item  
→ classified data

id	dim1	dim 2	...	dim n	class
1	69	23.5	...	4.3	A
2	54	27.4	...	2.7	C
3	81	22.4	...	5.2	B
4	72	31.5	...	1.5	C
5	69	25.4	...	4.8	A
...	...	...	...	...	...
m	78	15.7	...	5.1	C

## Main Topics

- Data preprocessing
  - Annotation
  - Missing Values
- Unsupervised Learning
  - Clustering
  - Data Visualization
- Supervised Learning
  - Classification (predict a class)
  - Regression (predict a value)
- Feature Engineering
- Feature Selection
- Dimension Reduction / Embedding

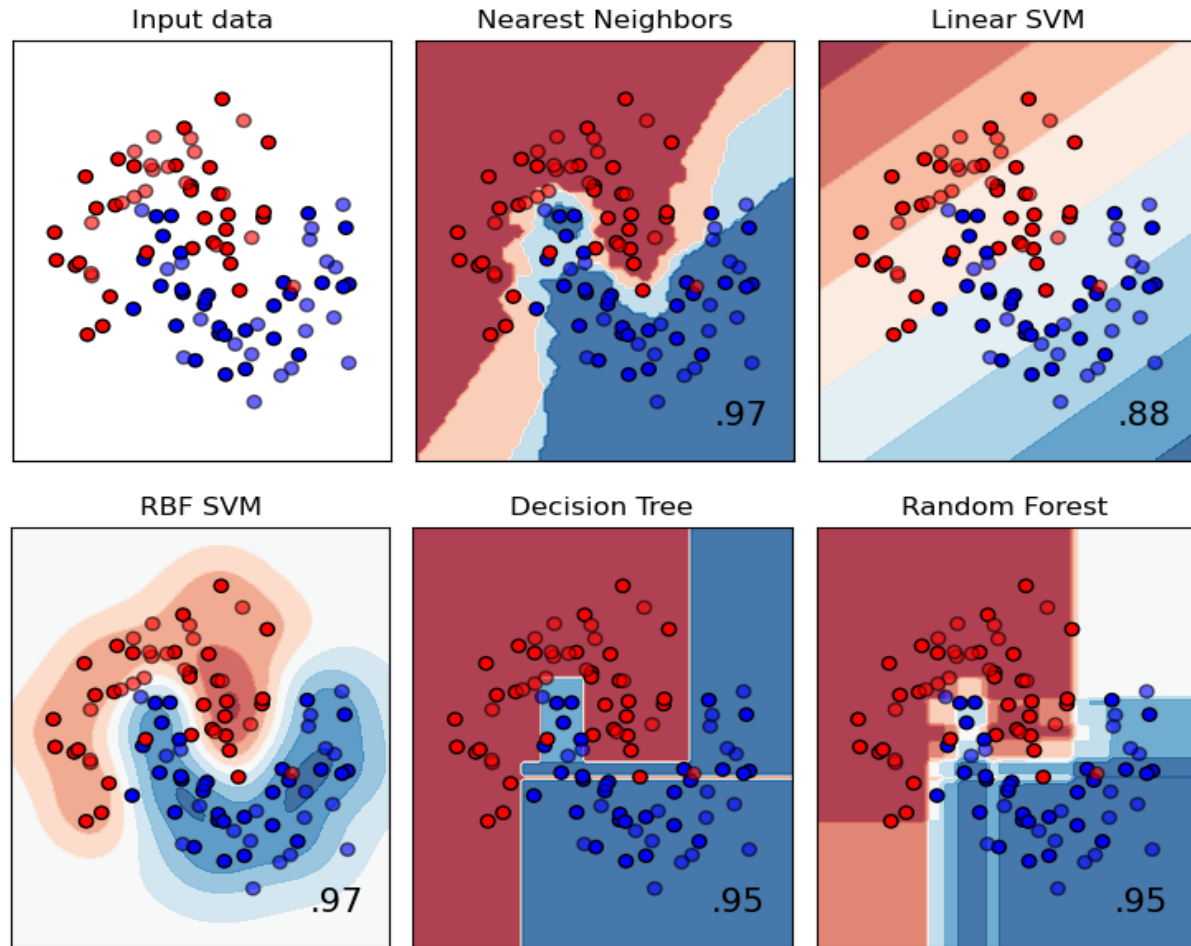
# Machine Learning

## Supervised Learning

- Train model on training data
  - paint feature space
- Evaluate model on test data
  - estimate class from position of sample in feature space
- Apply model on new data

## Supervised Learning Methods

- k-nearest neighbor (knn)
- Linear Regression
- Logistic Regression
- Support Vector Machines (SVM)
- Decision Trees / Random Forests
- Gaussian Process
- Naïve Bayes
- Neural Networks
- ...



Adapted from [https://scikit-learn.org/stable/auto\\_examples/classification/plot\\_classifier\\_comparison.html](https://scikit-learn.org/stable/auto_examples/classification/plot_classifier_comparison.html)

© 2007 - 2019, scikit-learn developers (BSD License).



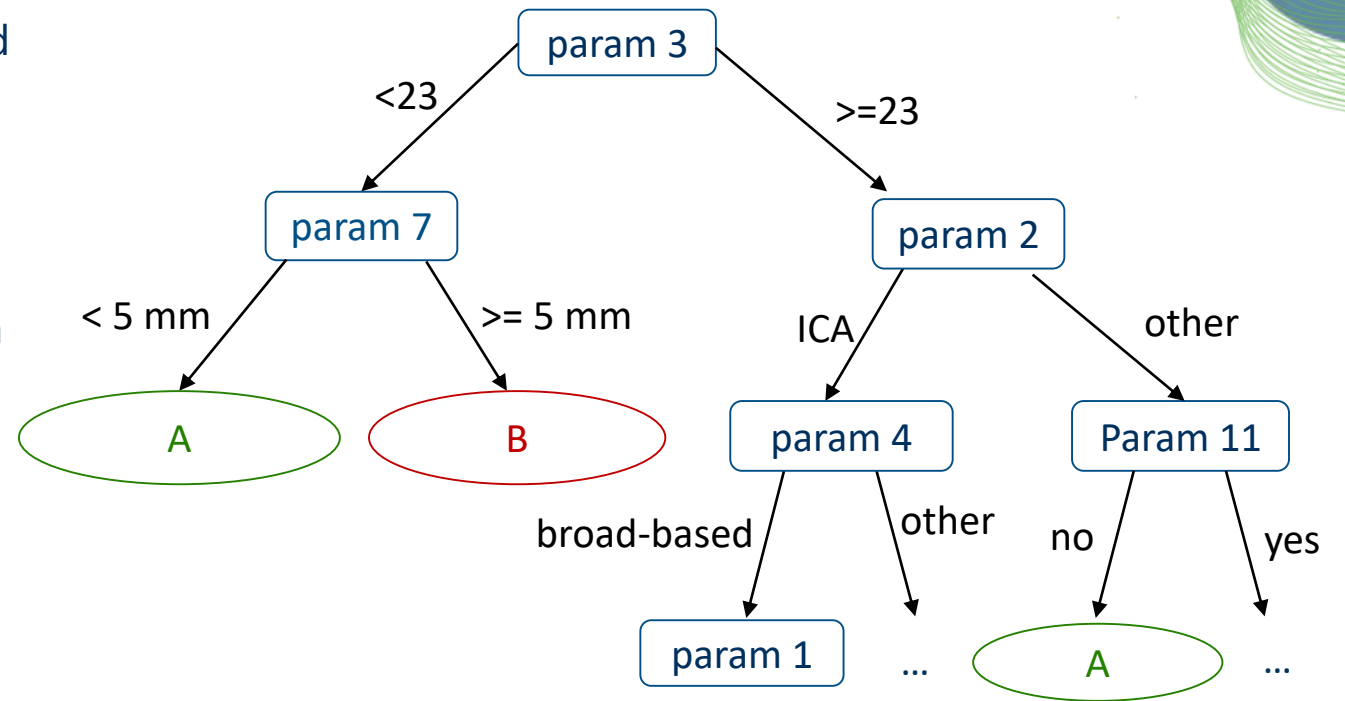
# Decision Tree

## Introduction

- machine learning algorithm
- data can be interval-scaled, categorical, or mixed
- classification: predict a class
- regression: predict a value
- shows good performance on tabular data (5-100 parameters, 50-1000 data points)
- model (tree) is computed based on training data

## Preparation

- Divide data in training data / test data
- Use 5-fold cross validation
  - 4/5 of data is training data
  - 1/5 of data is test data
  - Repeat 5 times
- Never train and test trained model on same data!



Decision Tree

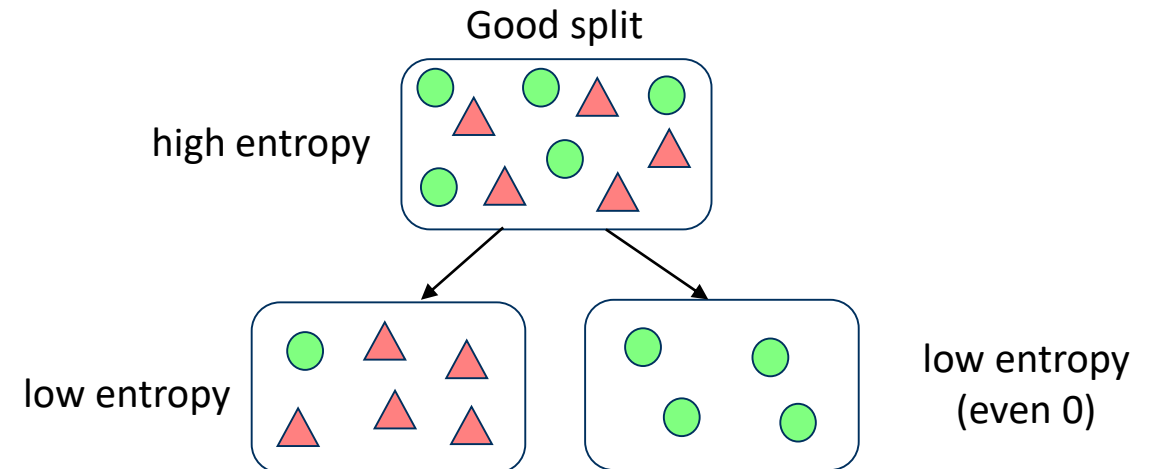
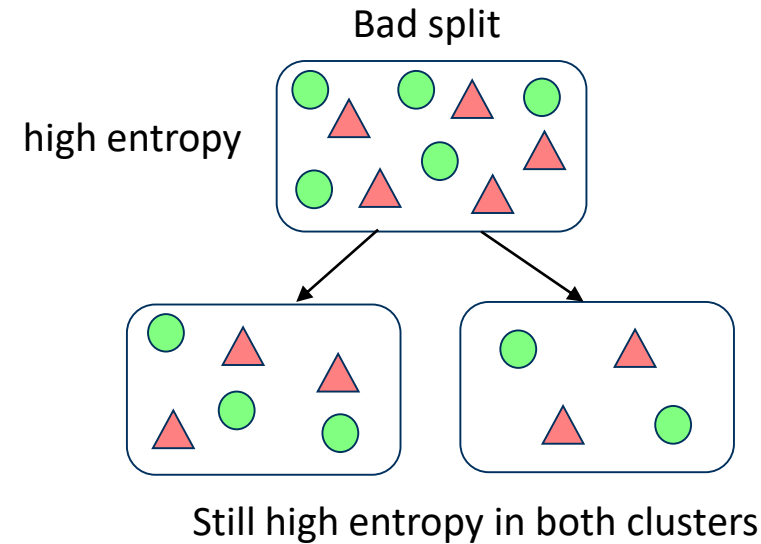
# Decision Tree - Training

## Training

- Start with complete data of training dataset
- For each step
  - choose parameter and threshold to minimize entropy in remaining clusters (leaves in tree)
  - Split cluster accordingly
- Entropy
  - measure for disorder
- Stopping criteria
  - Maximal depth reached (e.g. 10)
  - Minimal samples in leaf reached (e.g. 5)

## Classification / Application

- Apply tree on
  - test data (for testing) or
  - new data (for application)





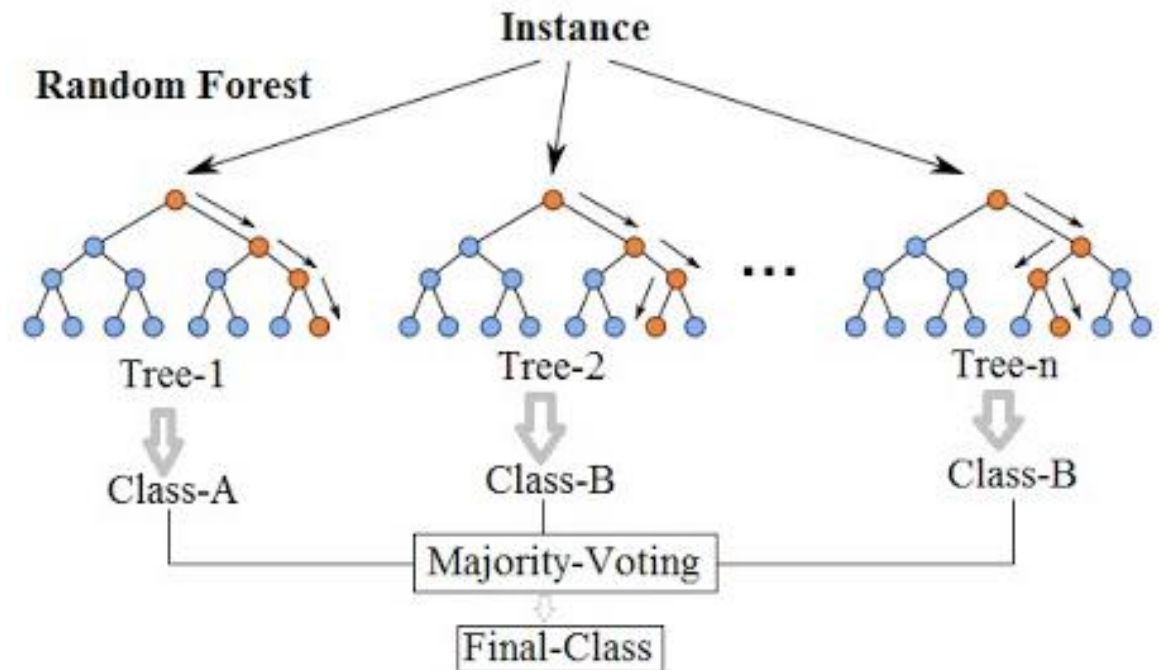
# Decision Trees and Random Forests

## Drawback of Decision Trees

- Problem
  - allow few levels → only few parameters are considered
  - allow many levels → overfitting
- Solution: Random Forests

## Random Forests

- Idea: train many decision trees with part of the data
- for each tree
  - use only part of the data items
  - use only part of the parameters
- Train  $n$  different trees ()
- Result:  $n$  slightly different decision trees
- Application: combine results using majority-voting



# Image Segmentation

# Image segmentation using thresholding

- Recap: Finding the right workflow towards a good segmentation takes time

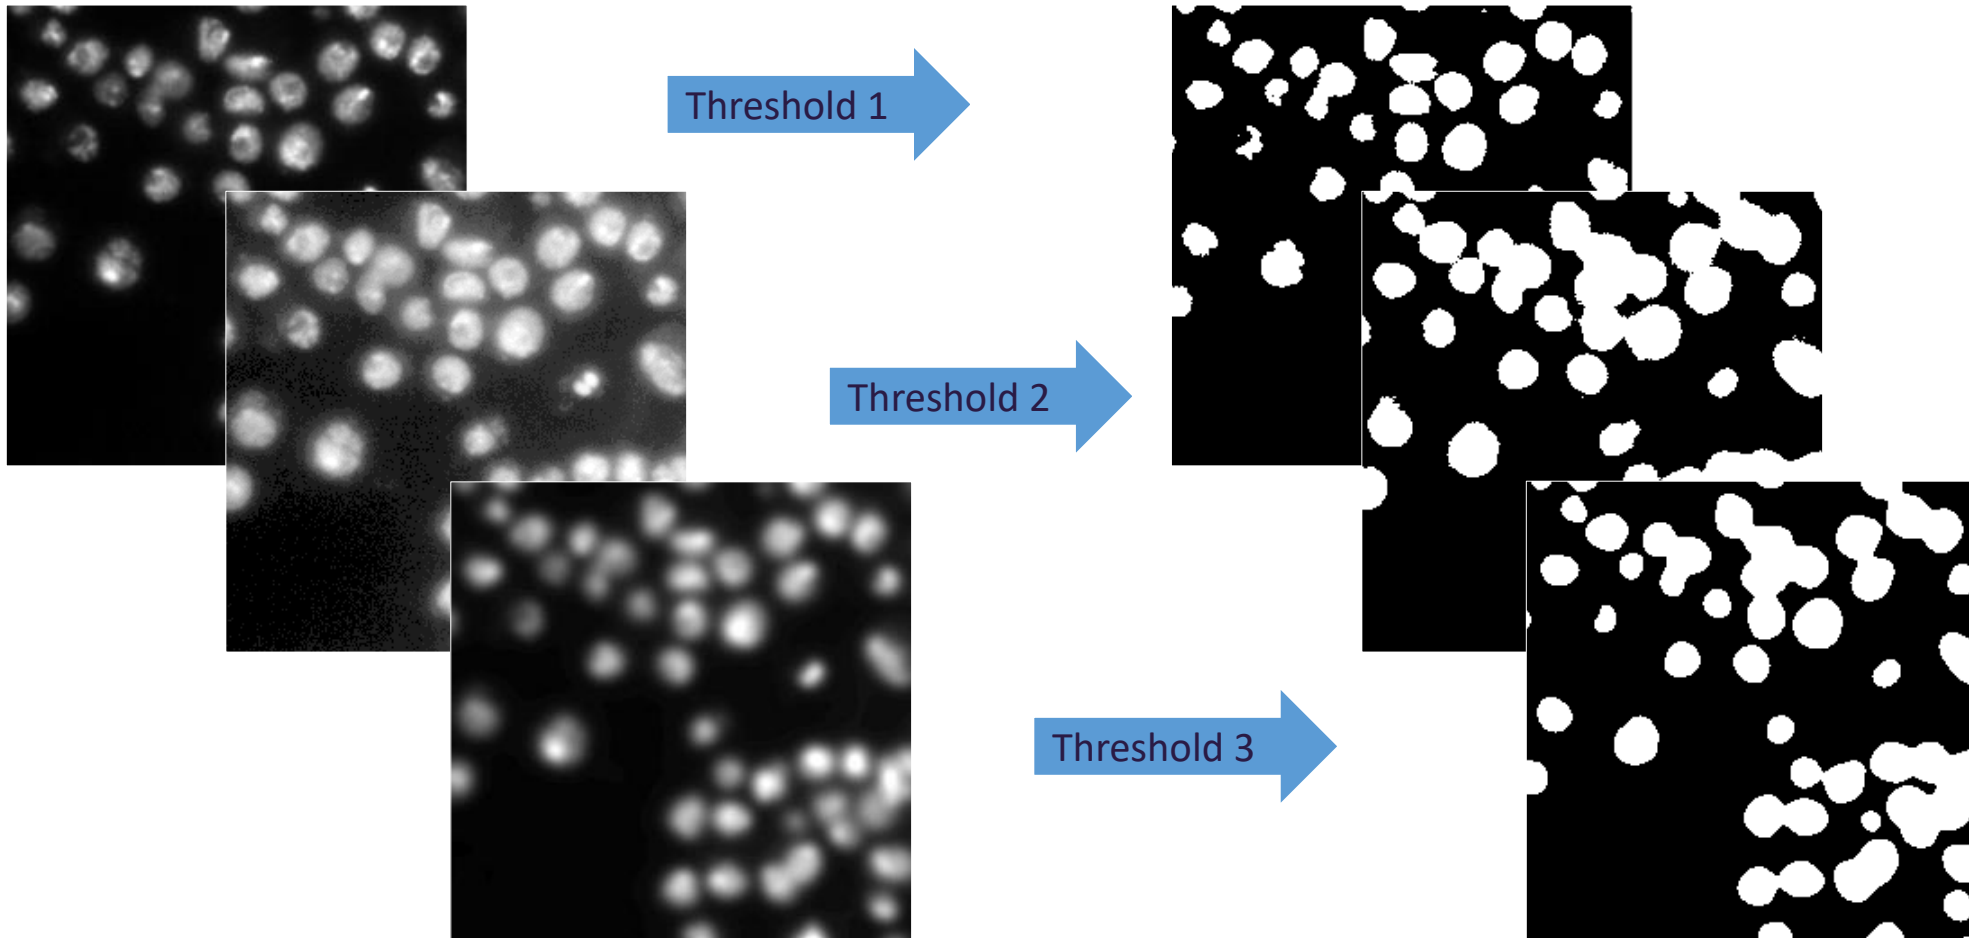


Image data source: [BBBC038v1](#), available from the Broad Bioimage Benchmark Collection (Caicedo et al., Nature Methods, 2019).

# Image segmentation using thresholding

- Recap: Recap: Combining images, e.g. using Difference of Gaussian (DoG)

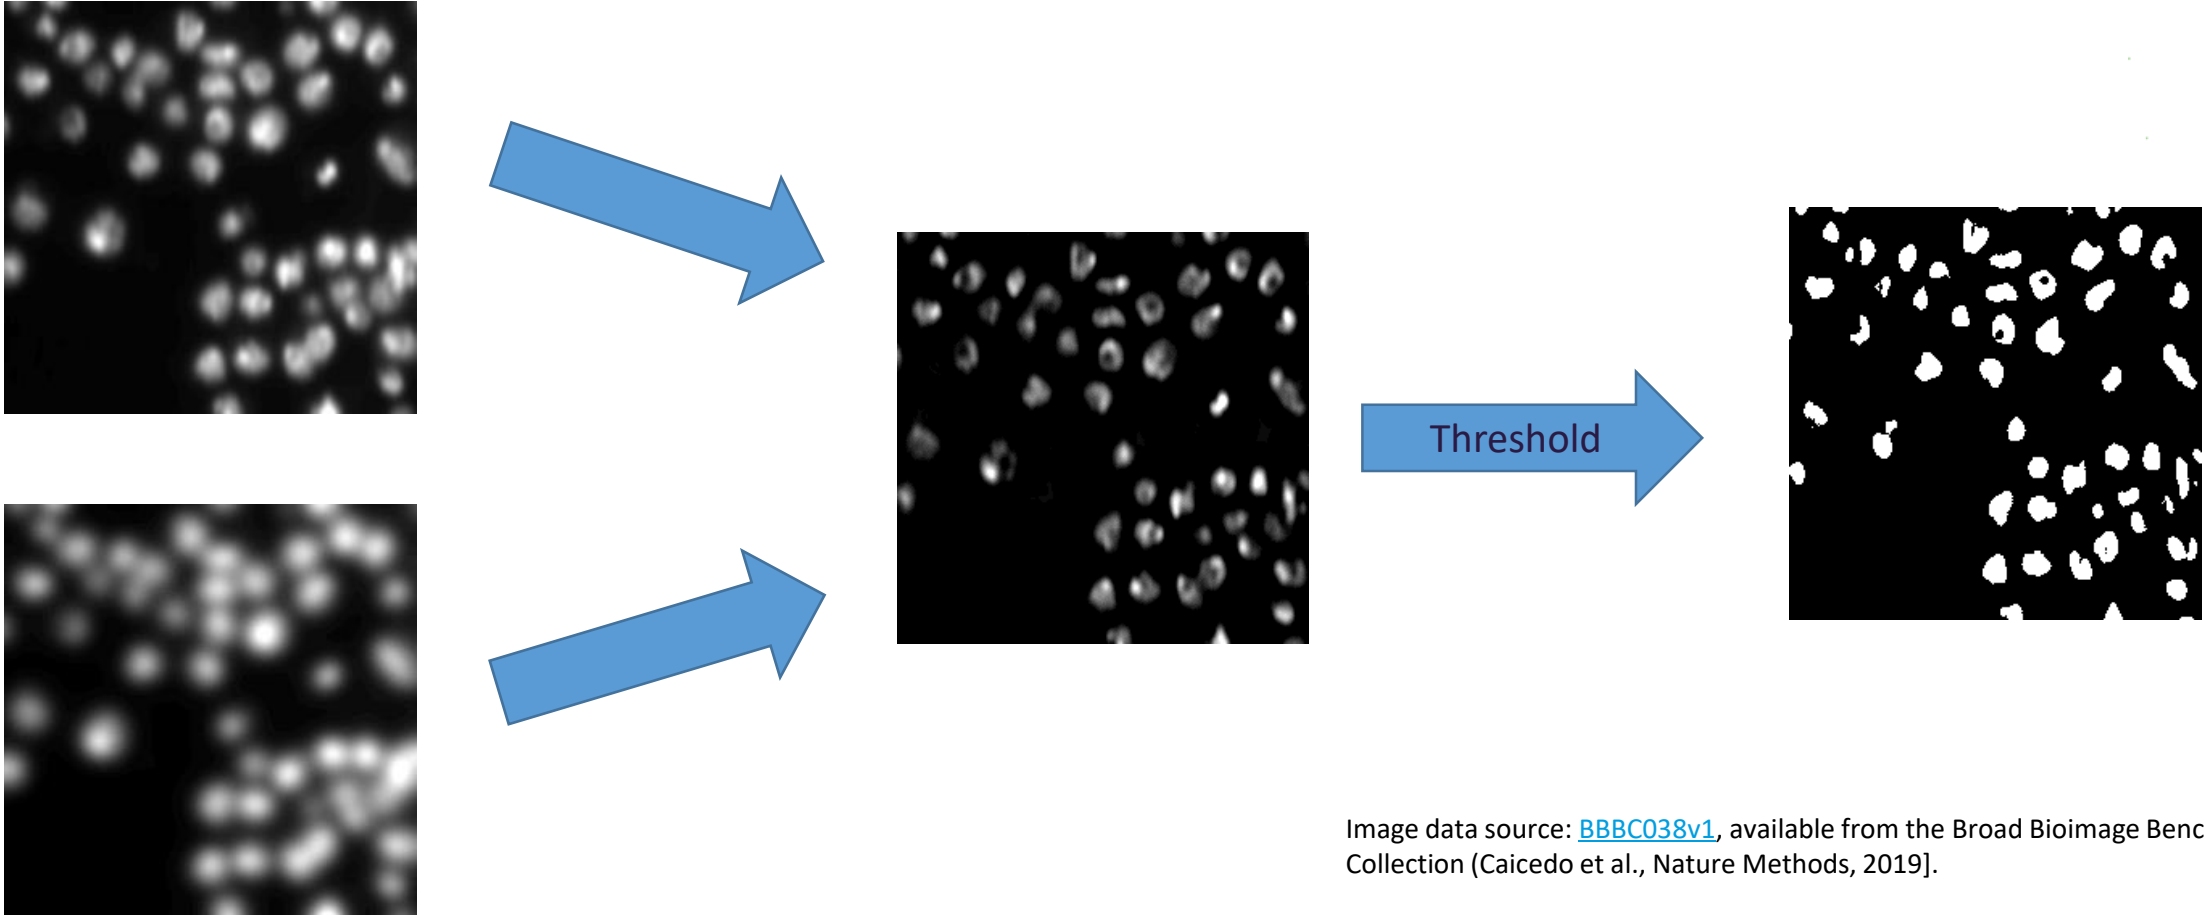


Image data source: [BBBC038v1](#), available from the Broad Bioimage Benchmark Collection (Caicedo et al., Nature Methods, 2019).

# Image segmentation using thresholding

- Might there be a technology for optimization which combination of images can be used to get the best segmentation result?

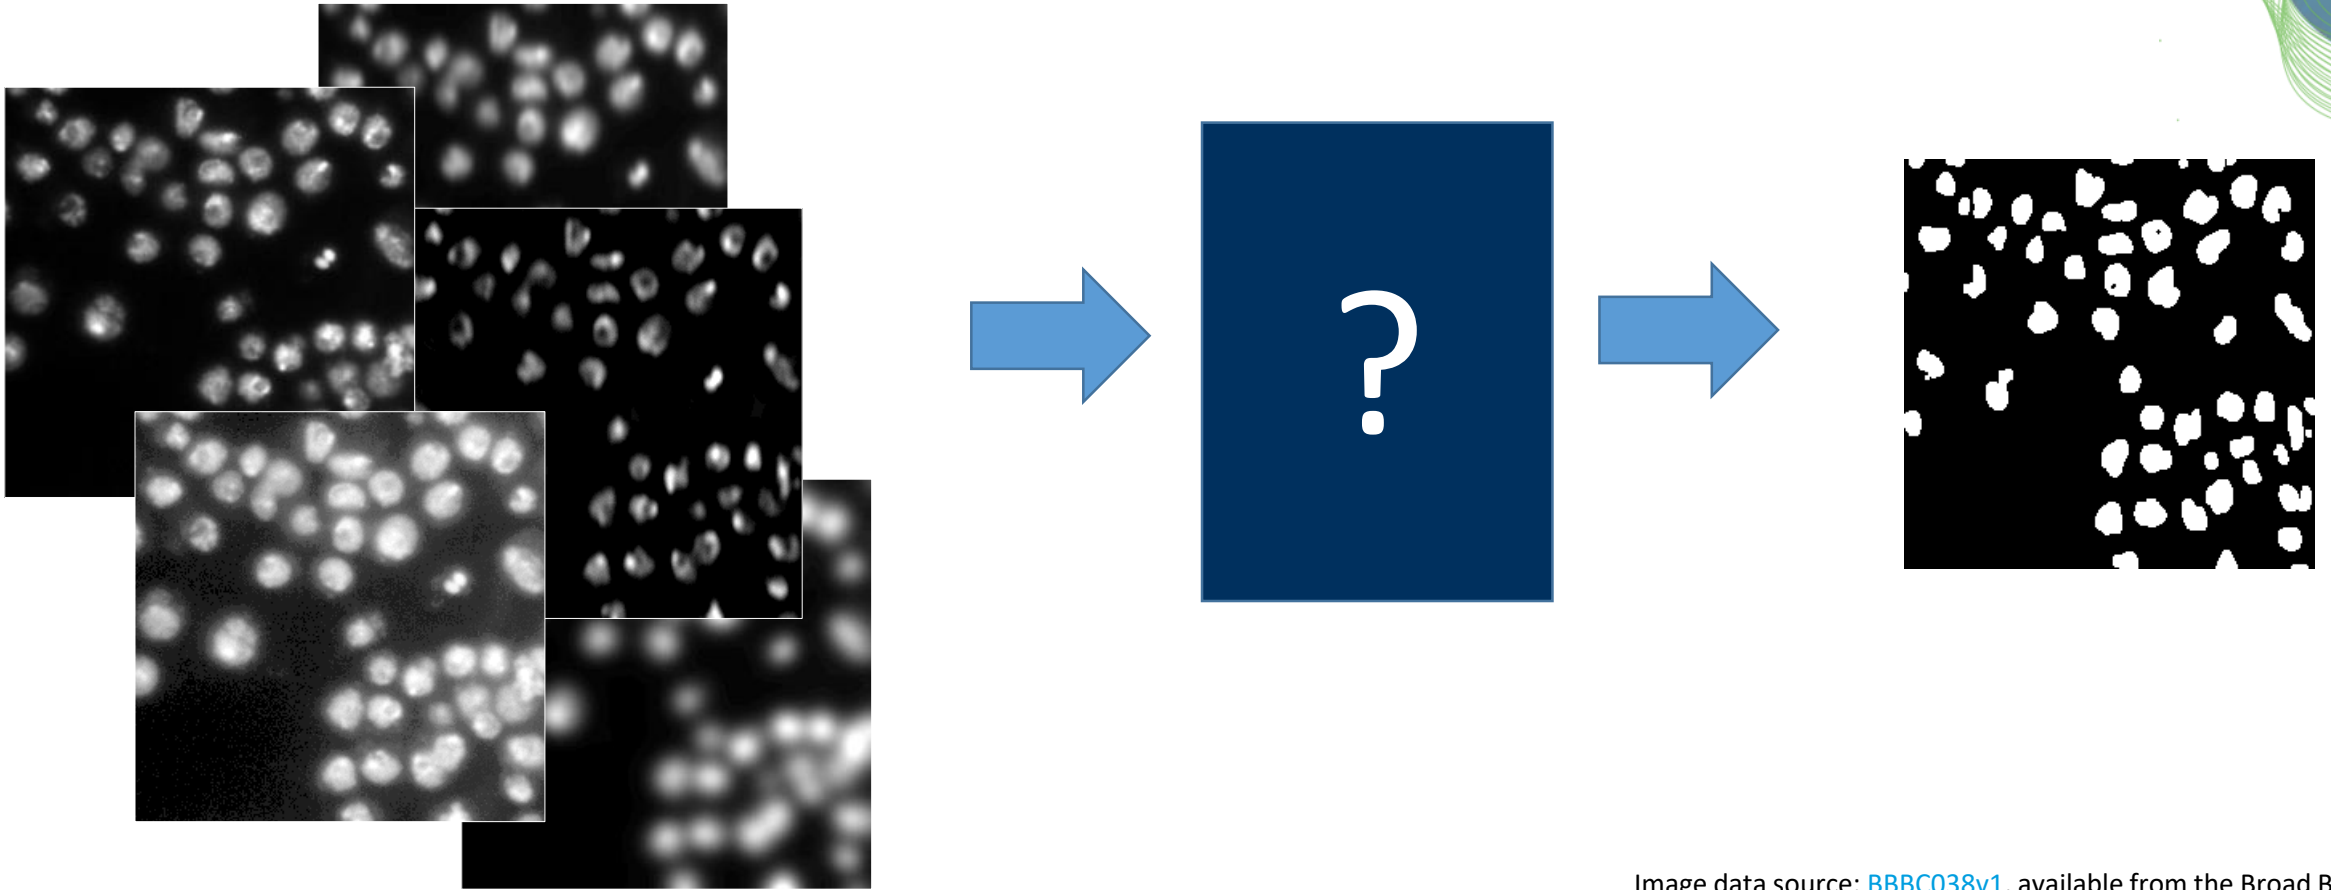


Image data source: [BBBC038v1](#), available from the Broad Bioimage Benchmark Collection (Caicedo et al., Nature Methods, 2019).

# Image segmentation using thresholding

- Might there be a technology for optimization which combination of images can be used to get the best segmentation result?

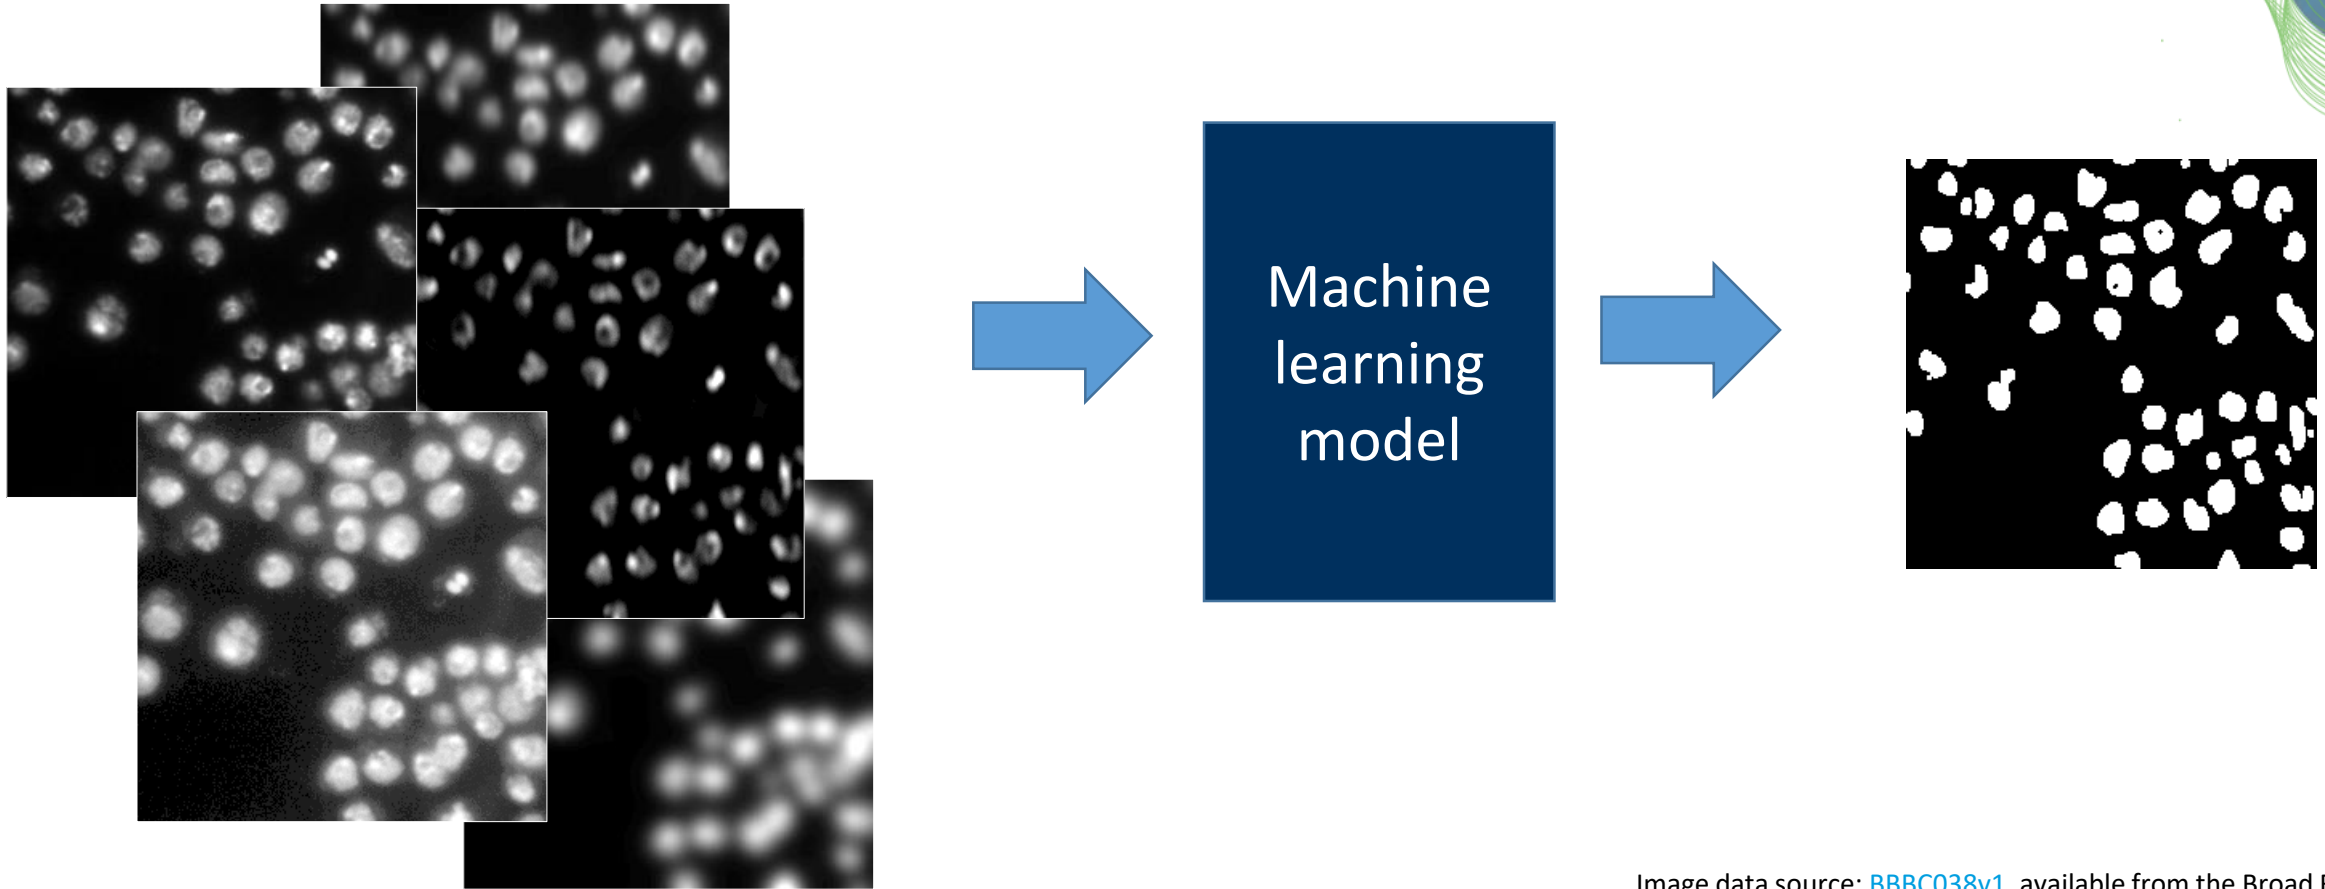
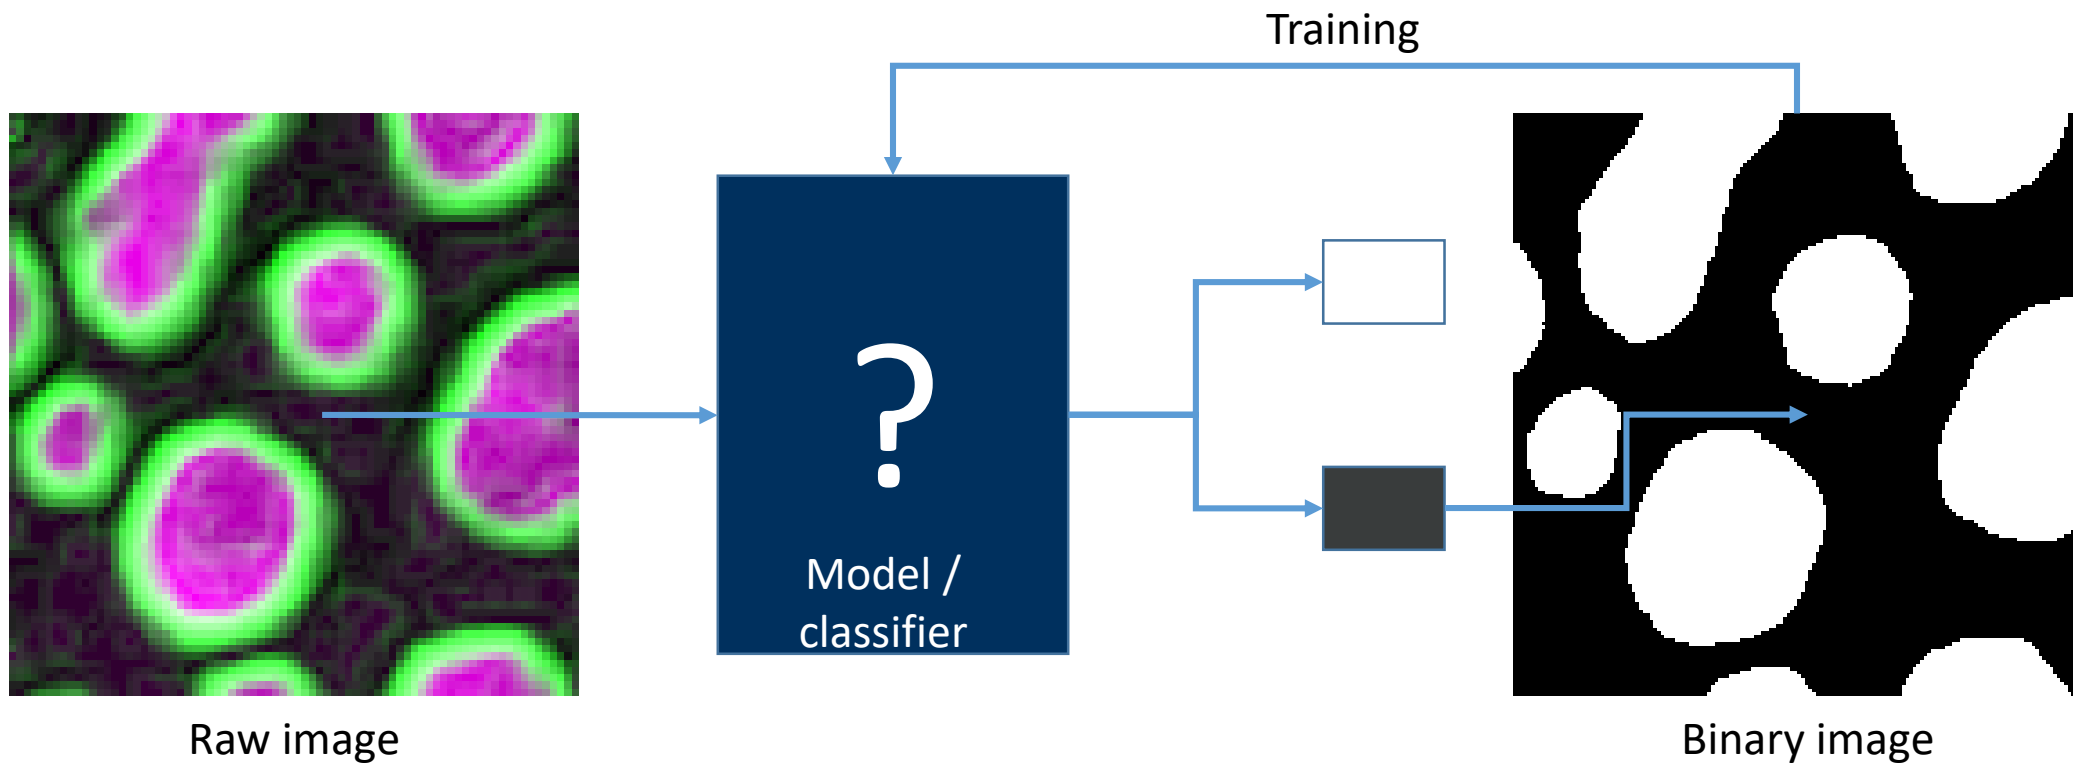


Image data source: [BBBC038v1](#), available from the Broad Bioimage Benchmark Collection (Caicedo et al., Nature Methods, 2019).



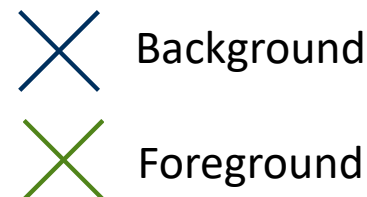
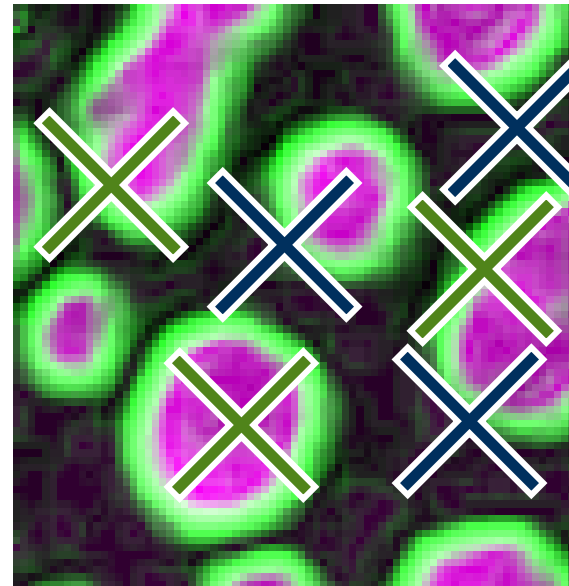
# Machine learning for image segmentation

- *Supervised* machine learning: We give the computer some ground truth to learn from
- The computer derives a *model* or a *classifier* which can judge if a pixel should be foreground (white) or background (black)
- Example: Binary classifier



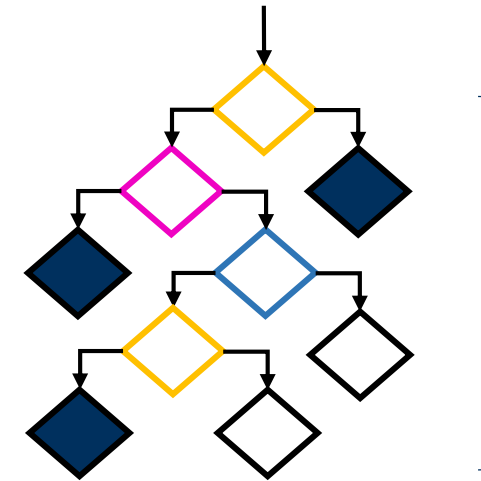
# Image segmentation using pixel classification

- Idea: use different features of a pixel to classify it to background or foreground
- Each pixel is considered separately
- Features:
  - Intensity/color of original pixel
  - Gaussian blur image
  - DoG image
  - LoG image
  - Hessian
- Features from different images
- For efficient processing, we randomly *sample* our dataset
- Create a dataset with pixel features vectors that belong to the background and the foreground
- Use machine learning (e.g. Random Forest) to classify each pixel



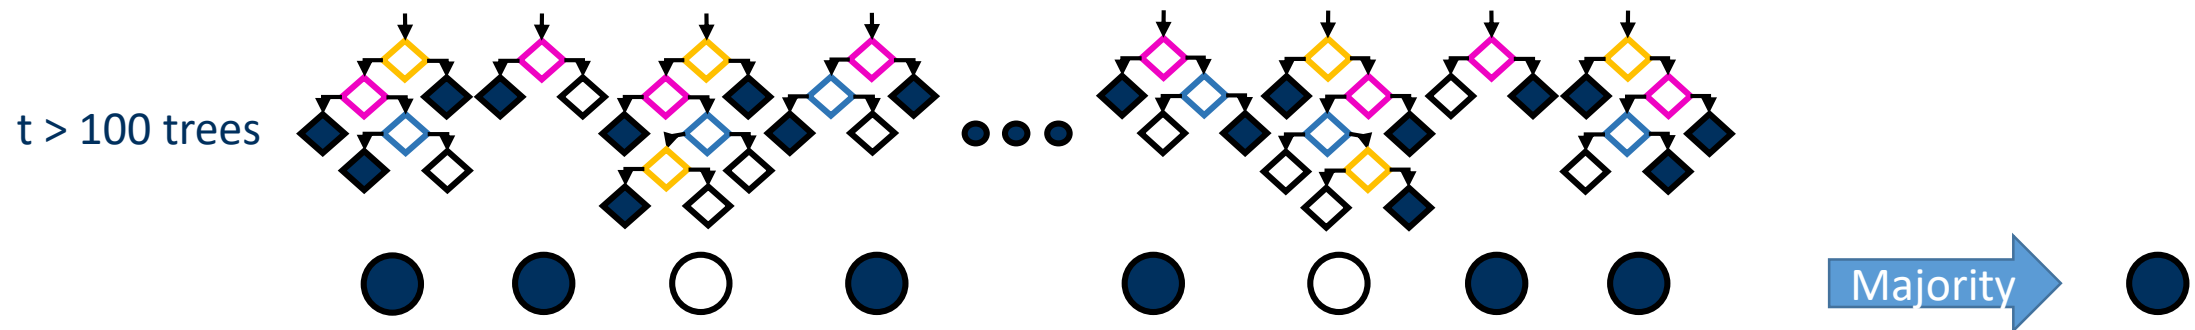
# Random Forest Pixel Classifier

Available features: > 20



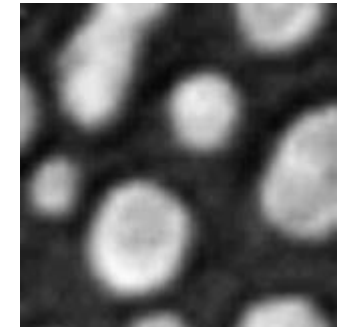
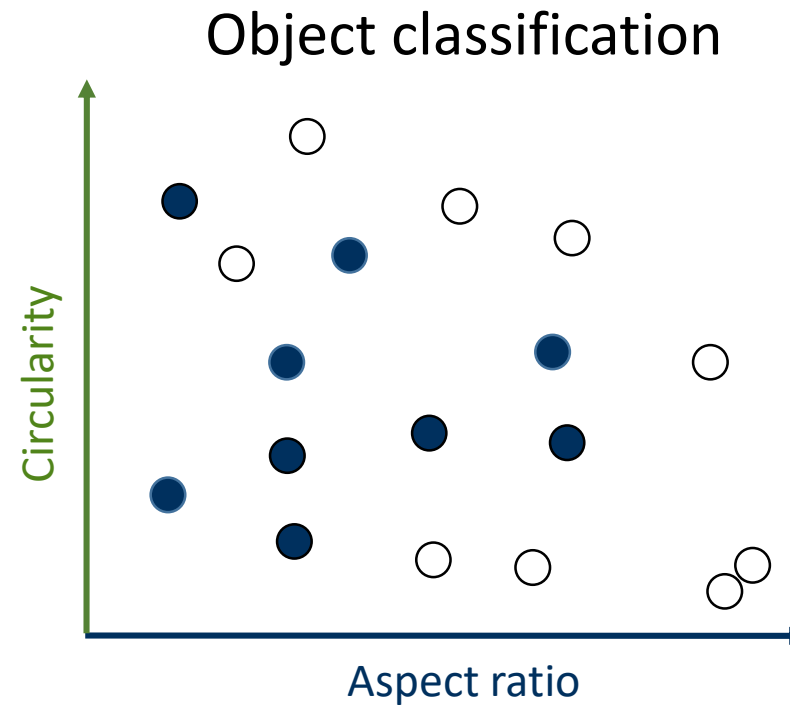
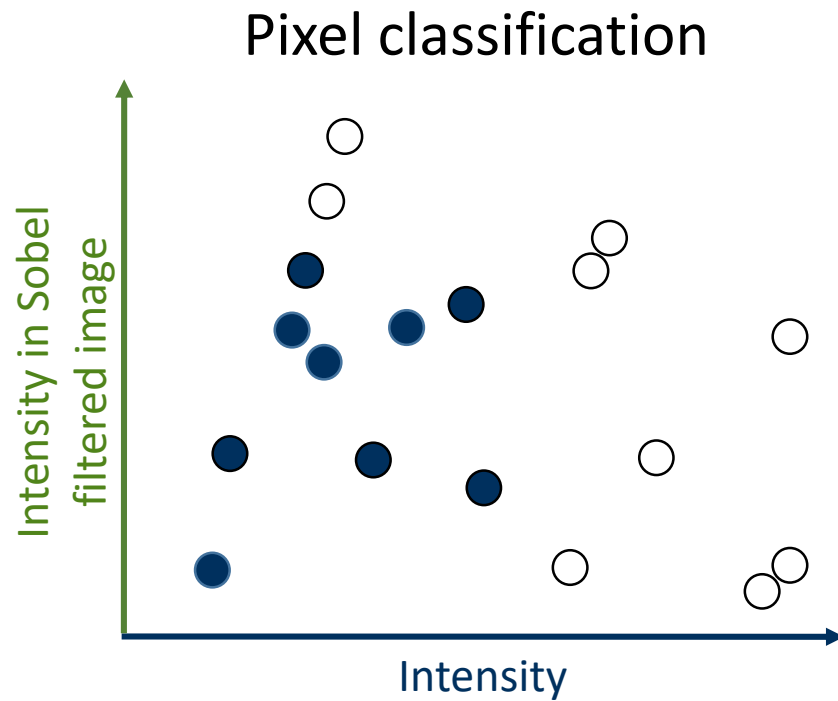
Depth: 4

- Train  $t$  trees on selected features and sampled pixels ->  $t$  different trees
- Combination of different tree decisions by max/mean voting



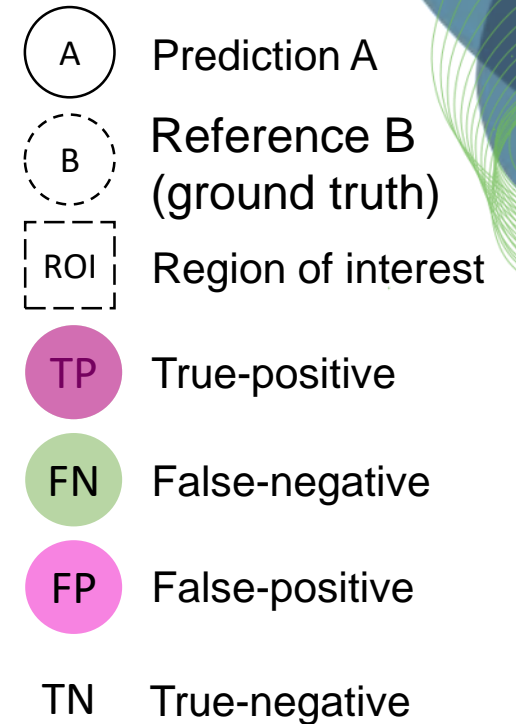
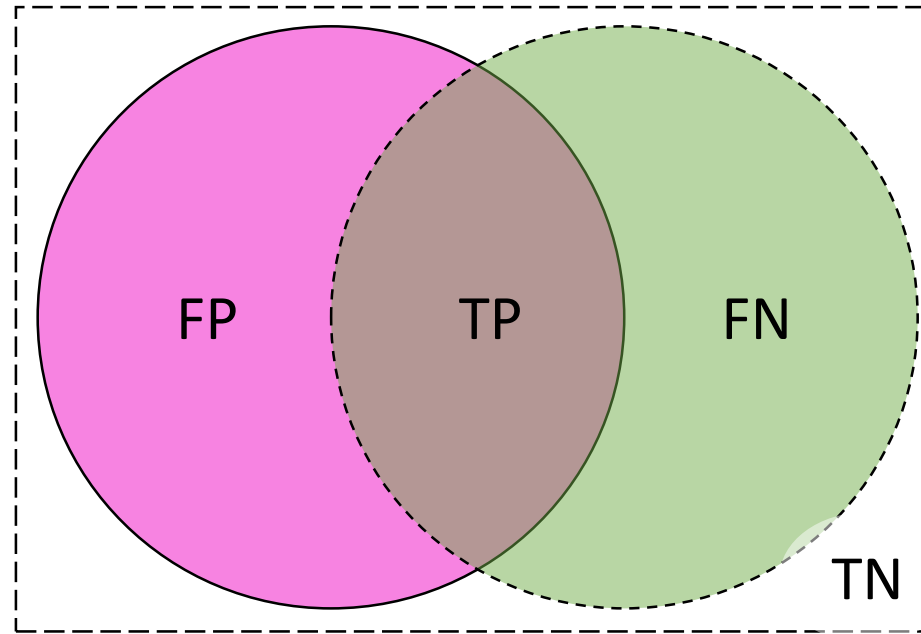
# Object classification

- Use object features instead of pixel features (e.g. size, aspect ratio, shape, circularity)
- The algorithms work the same



# Segmentation quality estimation

- In general
  - Define what's positive and what's negative.
  - Compare with a reference to figure out what was true and false
- Welcome to the Theory of Sets



Overlap  
(a.k.a. Jaccard index)

$$\frac{TP}{FP + TP + FN}$$

How much do A and B overlap?

Precision

$$\frac{TP}{TP + FP}$$

What fraction of points that were predicted as positives were really positive?

Recall  
(a.k.a. sensitivity)

$$\frac{TP}{TP + FN}$$

What fraction of positives points were predicted as positives?

# Model validation

- A good classifier is trained on a hand full of datasets and works on thousands similarly well.
- In order to assess that, we split the ground truth into two set
  - Training set (80% of the available data)
  - Test set (20% of the available data)

Typically done with hundreds or thousands of cells / images / objects / whatever.

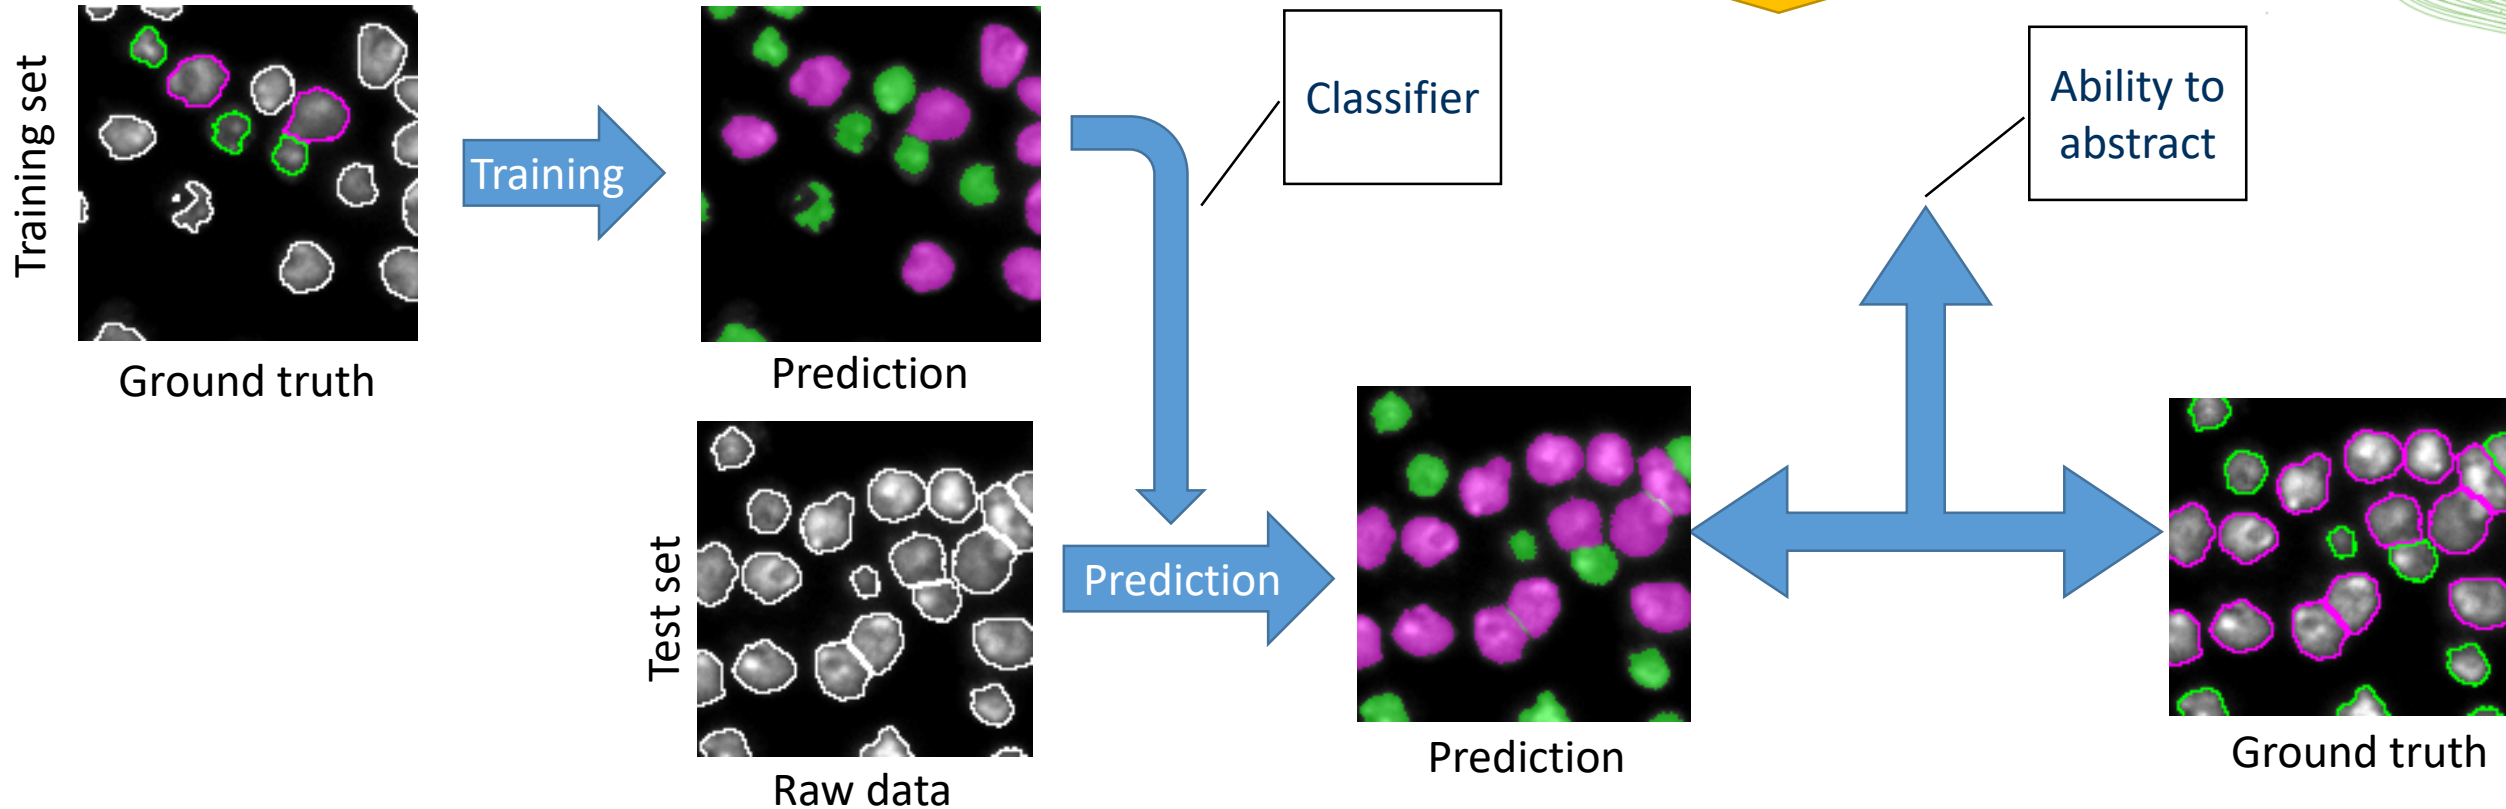


Image data source: [BBBC038v1](#), available from the Broad Bioimage Benchmark Collection (Caicedo et al., Nature Methods, 2019).



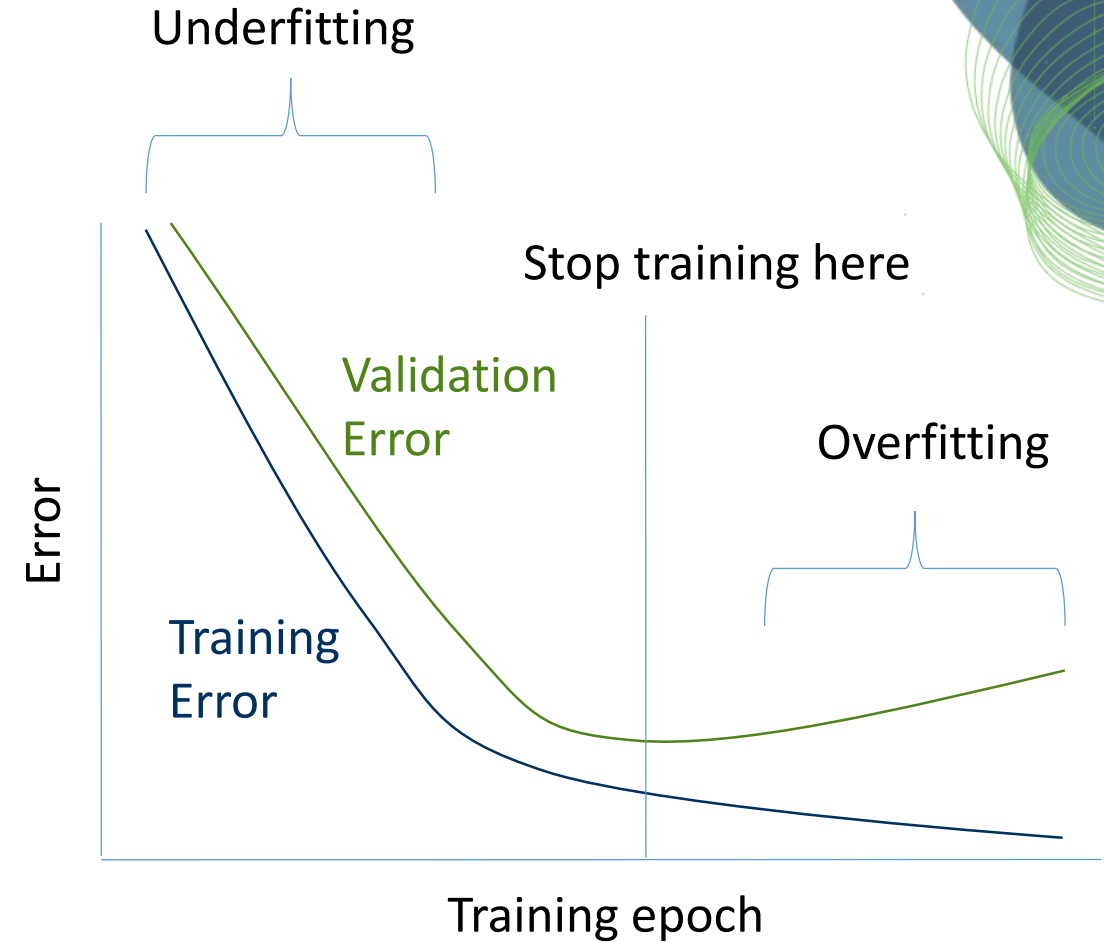
# Model validation

## Split data in

- Training dataset (80% of the data): used for training the model
- Validation dataset (10% of the data): after each iteration, see if the model overfits
- Test dataset (10% of the data): final evaluation after training is finished

## Training

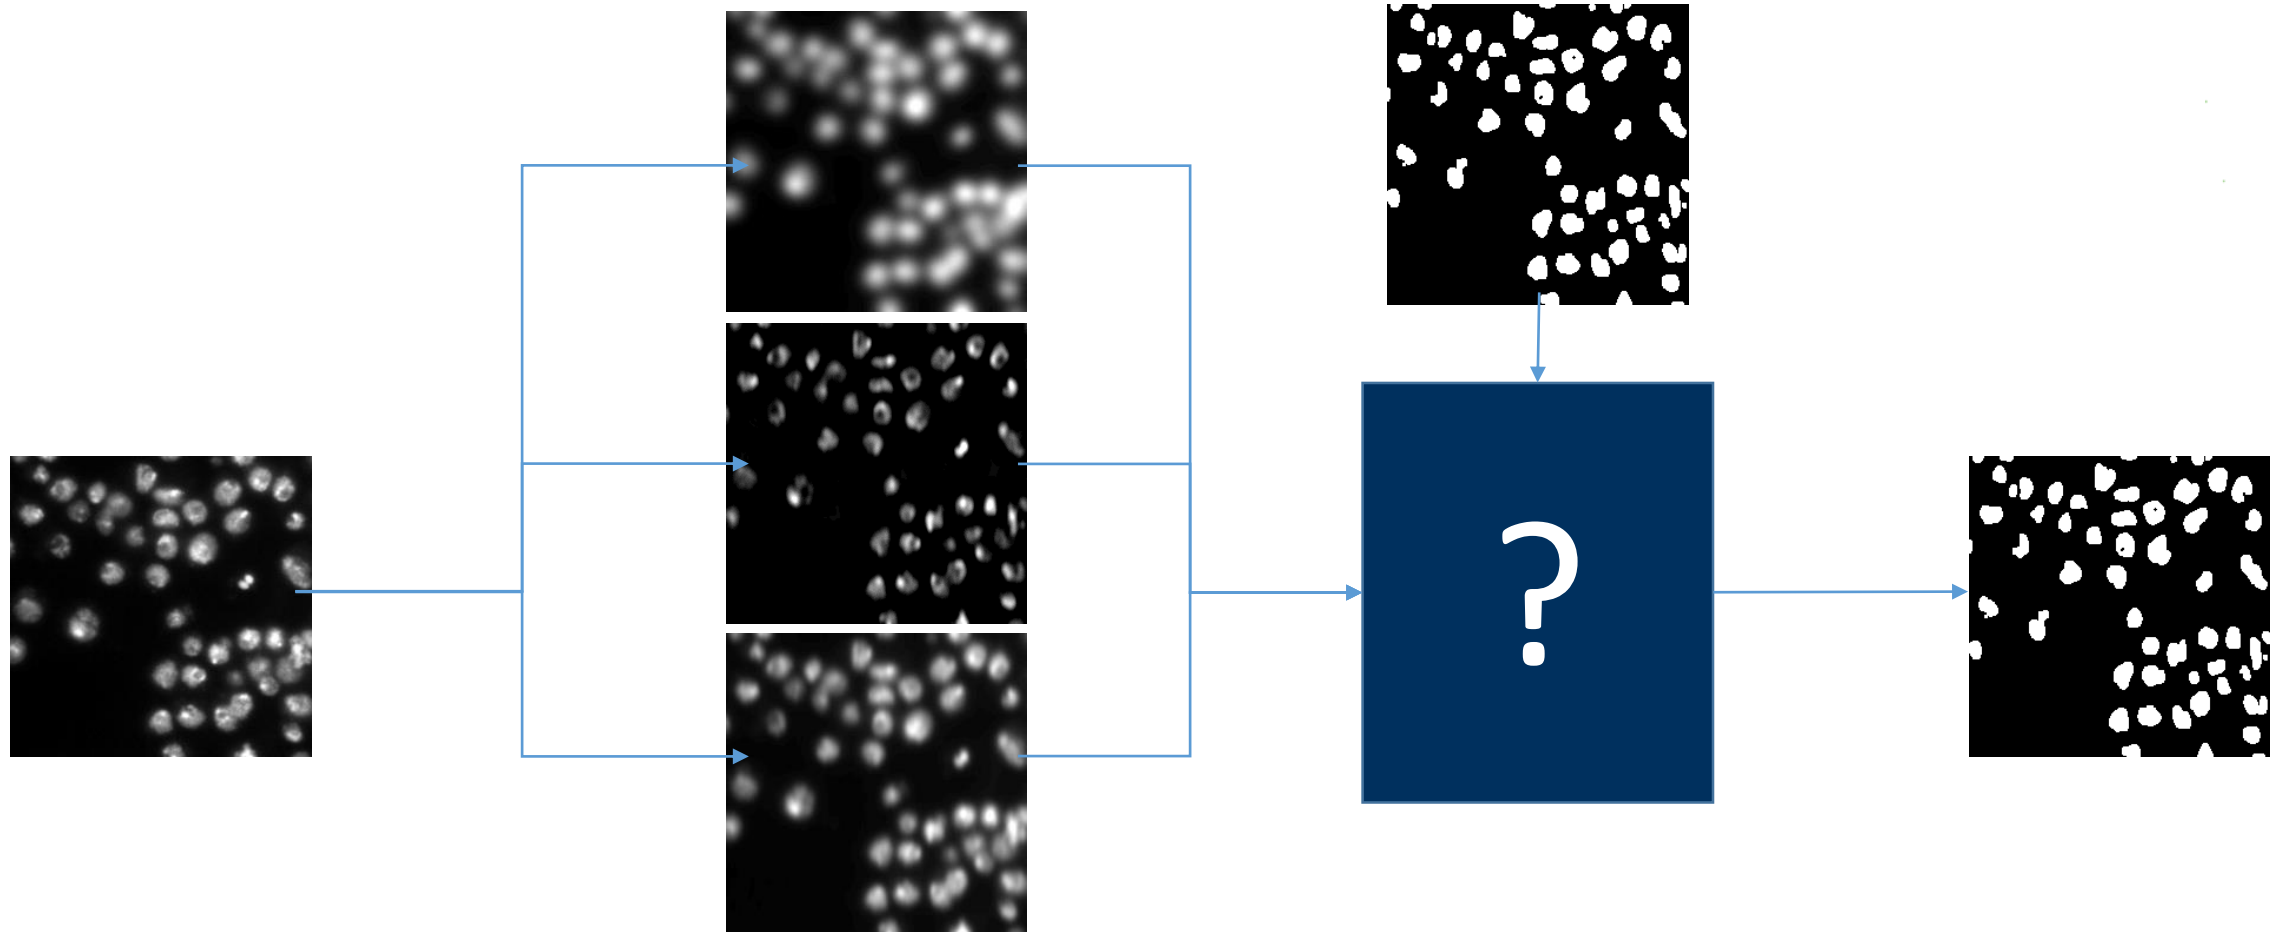
- Find spot with lowest validation error
- Avoid Underfitting: A model that is not trained long enough to capture the structure of the data
- Avoid Overfitting: A model that has been trained too long, has memorized the training data, but is not able to generalize on new data



<https://towardsdatascience.com/how-to-split-data-into-three-sets-train-validation-and-test-and-why-e50d22d3e54c>

# Outlook: Machine learning for image analysis

- In classical machine learning, we typically select features for training our classifier

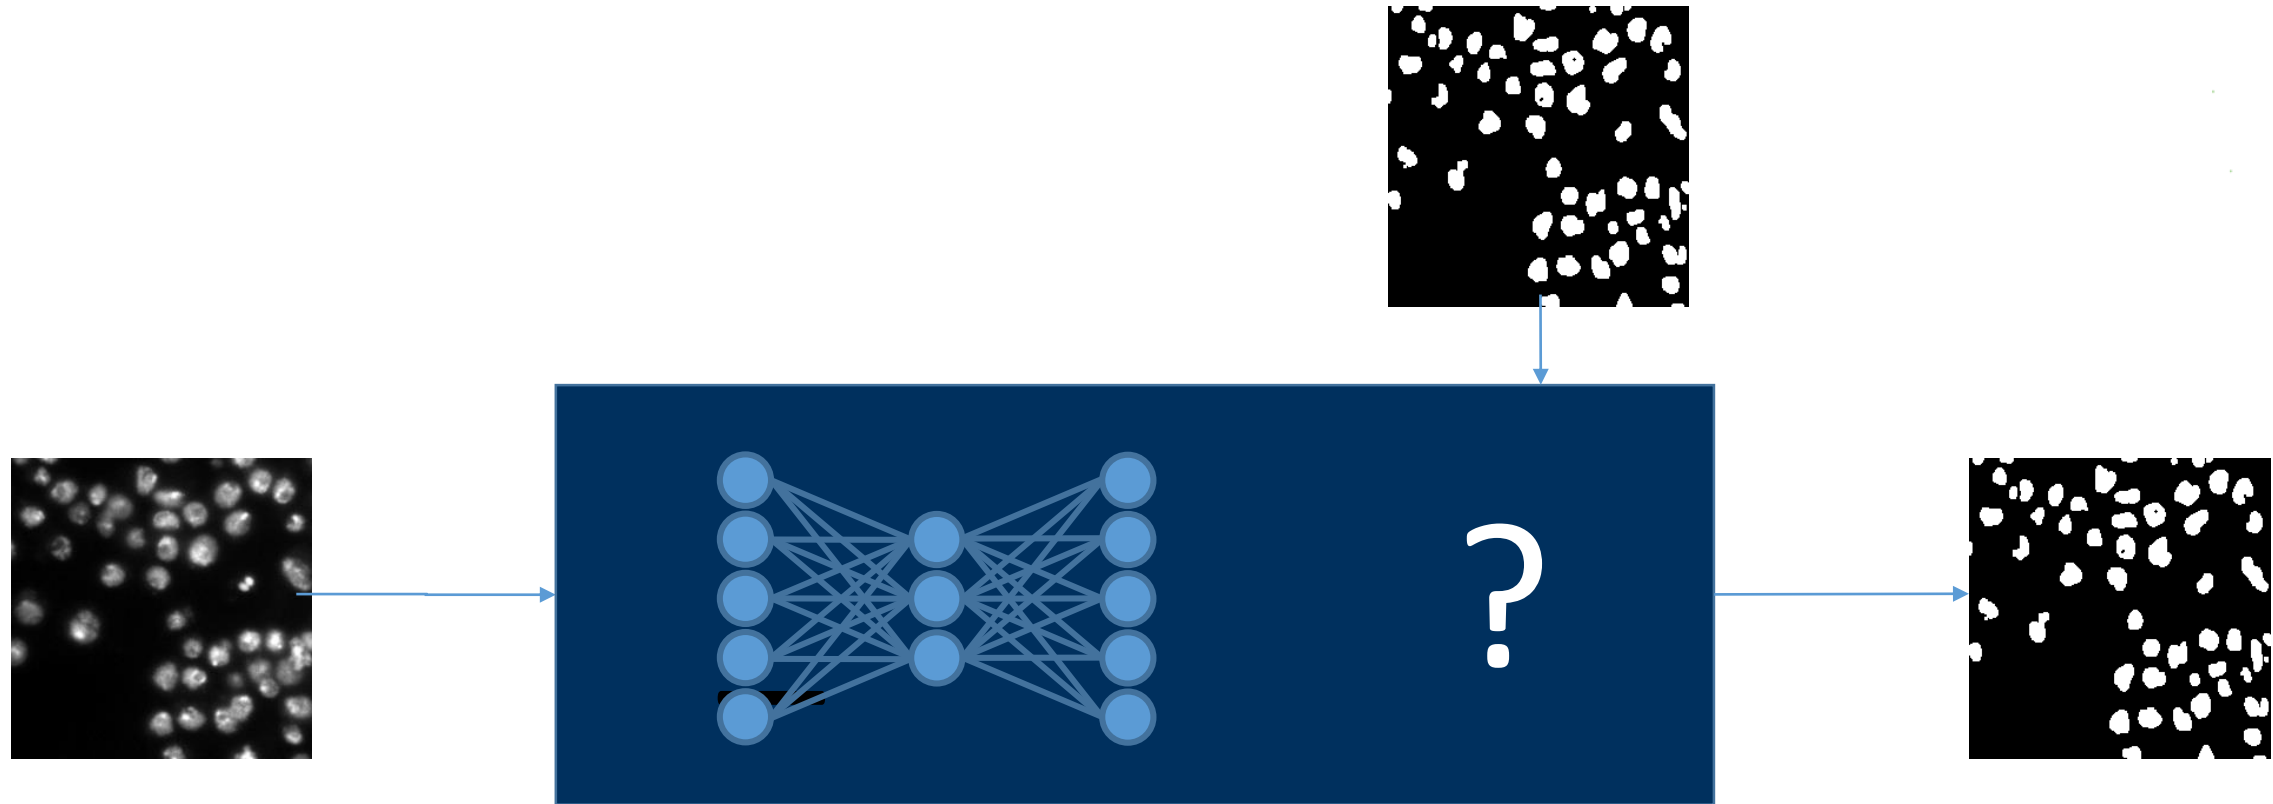


Convolutions

Image data source: [BBBC038v1](#), available from the Broad Bioimage Benchmark Collection (Caicedo et al., Nature Methods, 2019).

# Outlook: Deep learning for image analysis

- In deep learning, this is done automatically by the neural network



Convolutional neural networks

Image data source: [BBBC038v1](#), available from the Broad Bioimage Benchmark Collection (Caicedo et al., Nature Methods, 2019).

# Pixel classification using scikit-learn

With material from  
Robert Haase

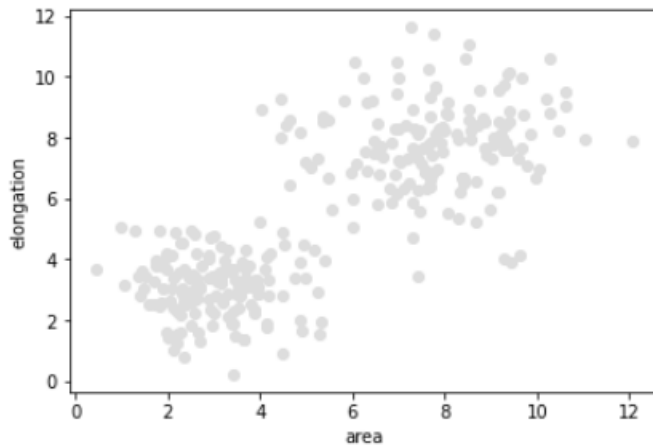
The scikit-learn logo is BSD3 licensed by the scikit-learn developers  
[https://commons.wikimedia.org/wiki/File:Scikit\\_learn\\_logo\\_small.svg](https://commons.wikimedia.org/wiki/File:Scikit_learn_logo_small.svg)

# Tabular object classification

- Classify objects starting from feature vectors (table columns)

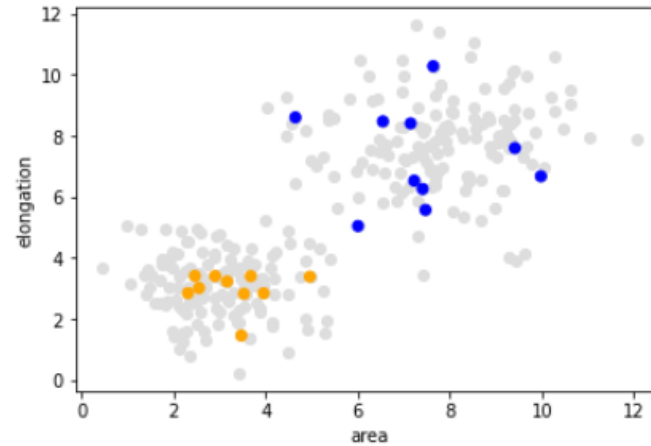
## Raw data

	area	elongation
0	3.950088	2.848643
1	4.955912	3.390093
2	7.469852	5.575289
3	2.544467	3.017479
4	3.465662	1.463756
5	3.156507	3.232181
6	9.978705	6.676372
7	6.001683	5.047063
8	2.457139	3.416050
9	3.672295	3.407462
10	9.413702	7.598608



## “Ground truth” annotation

annotation = [1, 1, 2, 1, 1, 1, 2, 2,

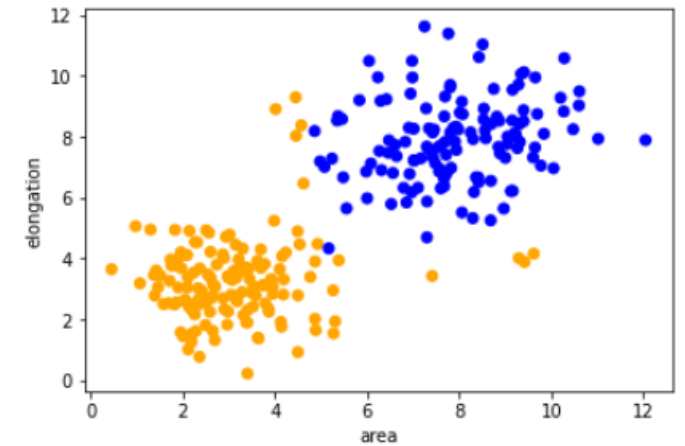


## Classifier training

```
classifier = RandomForestClassifier()  
classifier.fit(train_data, train_annotation)
```

## Classifier prediction

```
result = classifier.predict(validation_data)
```



[https://github.com/BiAPoL/Bio-image\\_Analysis\\_with\\_Python/blob/main/09\\_machine\\_learning/01\\_supervised\\_machine\\_learning.ipynb](https://github.com/BiAPoL/Bio-image_Analysis_with_Python/blob/main/09_machine_learning/01_supervised_machine_learning.ipynb)

# Interactive pixel classification

- Prepare an empty layer for annotations and keep a **reference**

```
labels = viewer.add_labels(  
    np.zeros(image.shape).astype(int))
```

- Read annotations

```
manual_annotations = labels.data
```

```
from skimage.io import imshow  
imshow>manual_annotations,  
    vmin=0, vmax=2)
```

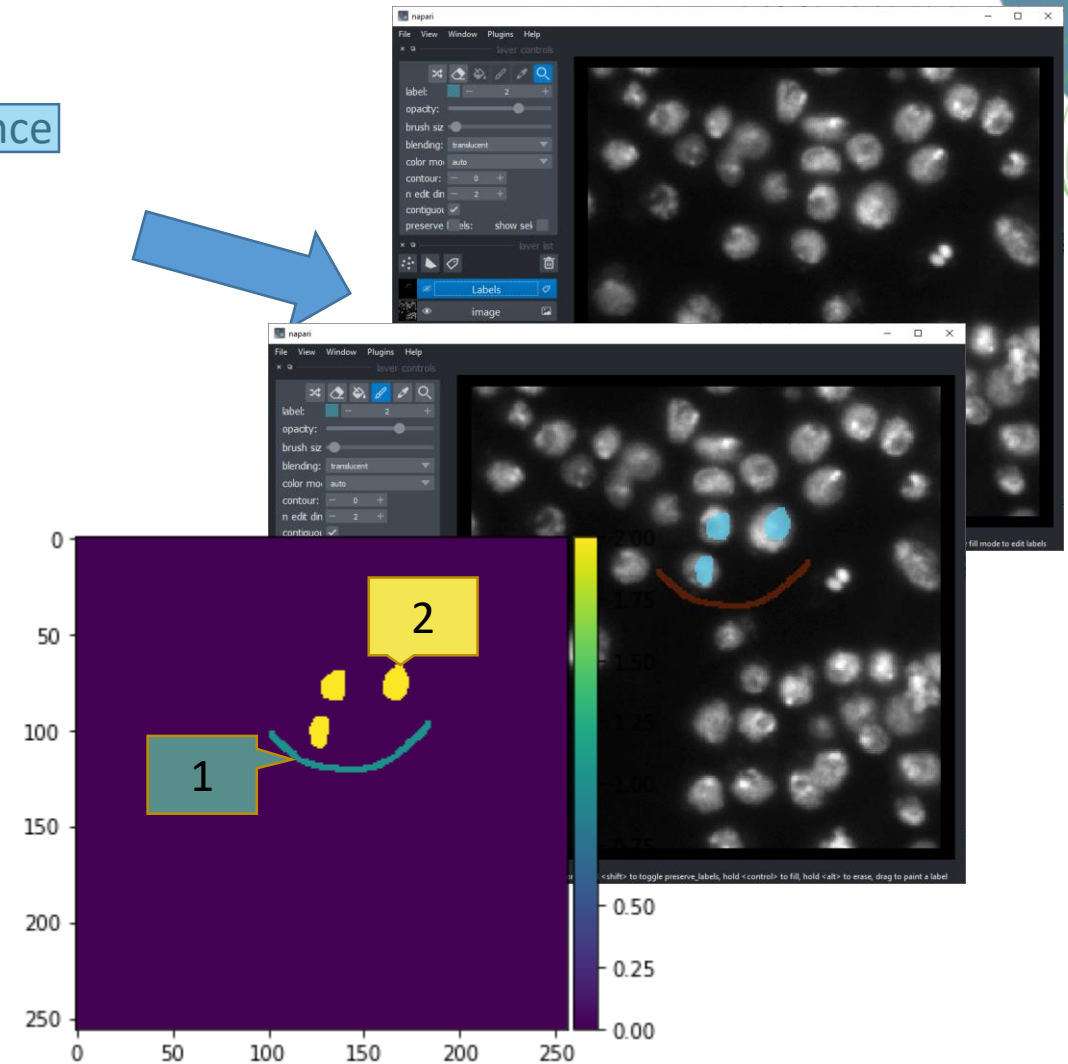


Image data source: [BBC038v1](https://github.com/BBC038v1), available from the Broad Bioimage Benchmark Collection (Caicedo et al., Nature Methods, 2019).

[https://github.com/BiAPoL/Bio-image\\_Analysis\\_with\\_Python/blob/main/09\\_machine\\_learning/02\\_scikit\\_learn\\_random\\_forest\\_pixel\\_classifier.ipynb](https://github.com/BiAPoL/Bio-image_Analysis_with_Python/blob/main/09_machine_learning/02_scikit_learn_random_forest_pixel_classifier.ipynb)



# Interactive pixel classification

- Pixel classification using scikit-learn
  - Expects one-dimensional arrays for
    - every feature individually
    - ground truth

```
# train classifier
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
classifier = RandomForestClassifier(max_depth=2, random_state=0)
```

```
classifier.fit(X, y)
```

Image data

Ground truth /  
annotation

Image data

```
y_ = classifier.predict(X)
```

prediction

[https://github.com/BiAPoL/Bio-image\\_Analysis\\_with\\_Python/blob/main/09\\_machine\\_learning/02\\_scikit\\_learn\\_random\\_forest\\_pixel\\_classifier.ipynb](https://github.com/BiAPoL/Bio-image_Analysis_with_Python/blob/main/09_machine_learning/02_scikit_learn_random_forest_pixel_classifier.ipynb)

# Interactive pixel classification

- Pixel classification using scikit-learn
  - Expects one-dimensional arrays for
    - every feature individually
    - ground truth

```
# for training, we need to generate features
```

```
feature_stack = generate_feature_stack(image)
```

```
X, y = format_data(feature_stack, manual_annotations)
```

```
# train classifier
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
classifier = RandomForestClassifier(max_depth=2, random_state=0)
```

```
classifier.fit(X, y)
```

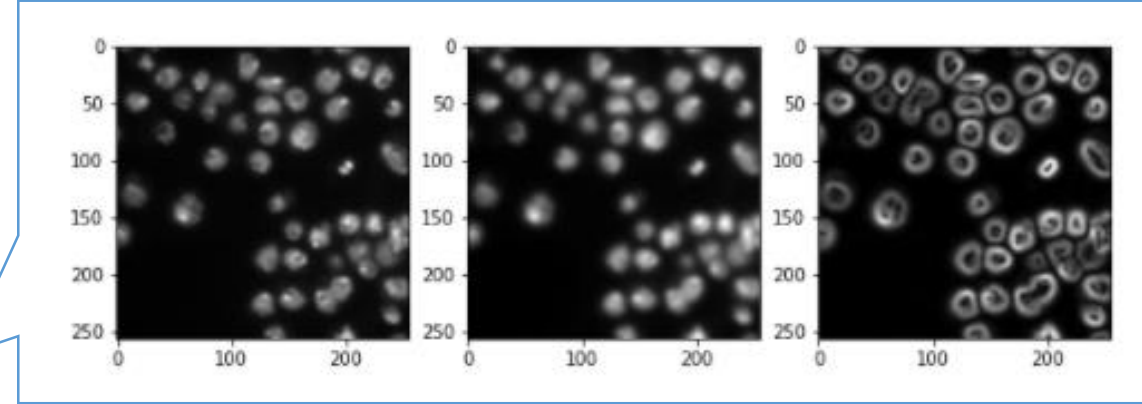


Image data source: [BBBC038v1](https://bbbc038v1), available from the Broad Bioimage Benchmark Collection (Caicedo et al., Nature Methods, 2019).

[https://github.com/BiAPoL/Bio-image\\_Analysis\\_with\\_Python/blob/main/09\\_machine\\_learning/02\\_scikit\\_learn\\_random\\_forest\\_pixel\\_classifier.ipynb](https://github.com/BiAPoL/Bio-image_Analysis_with_Python/blob/main/09_machine_learning/02_scikit_learn_random_forest_pixel_classifier.ipynb)

# Interactive pixel classification

- Pixel classification using scikit-learn

```
# process the whole image and show result
result_1d =
classifier.predict(feature_stack.T)
result_2d = result_1d.reshape(image.shape)

viewer.add_labels(result_2d)
```

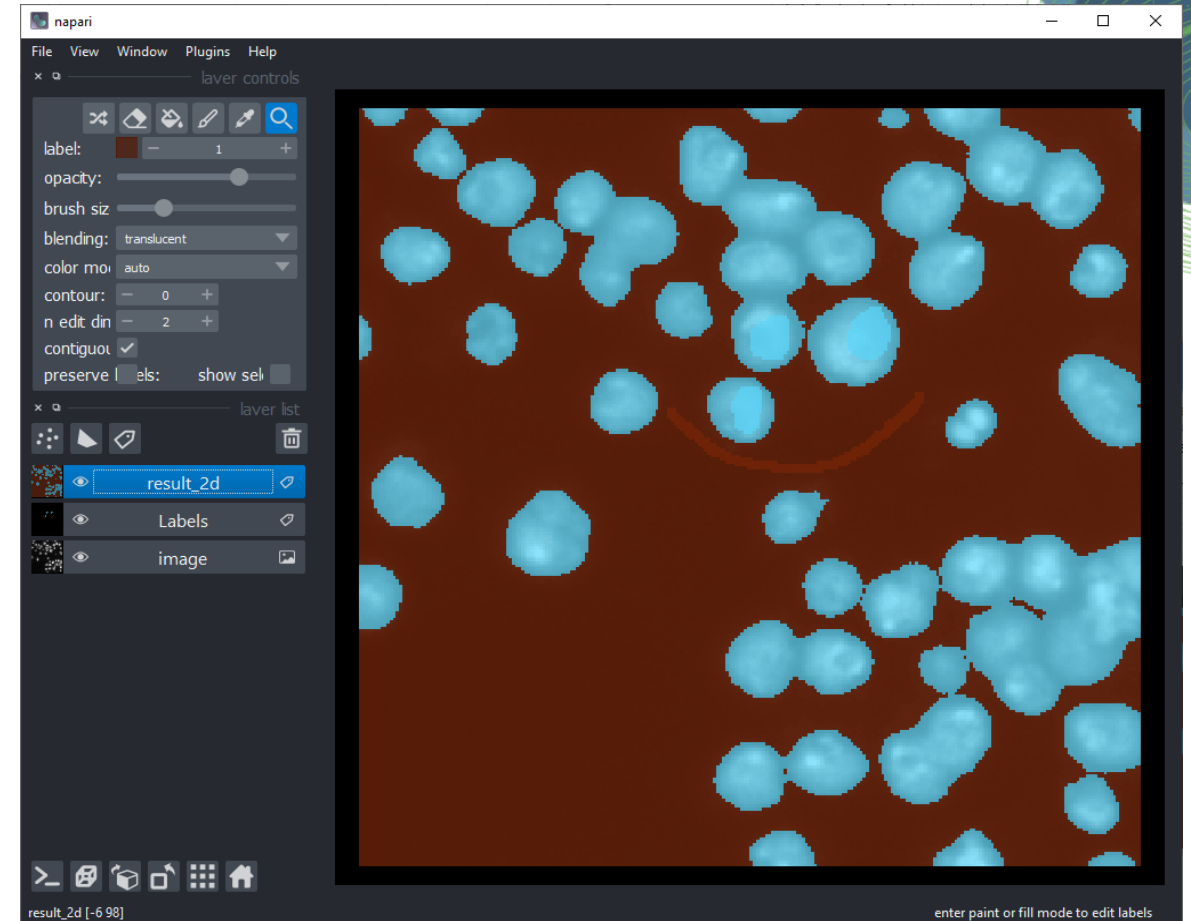
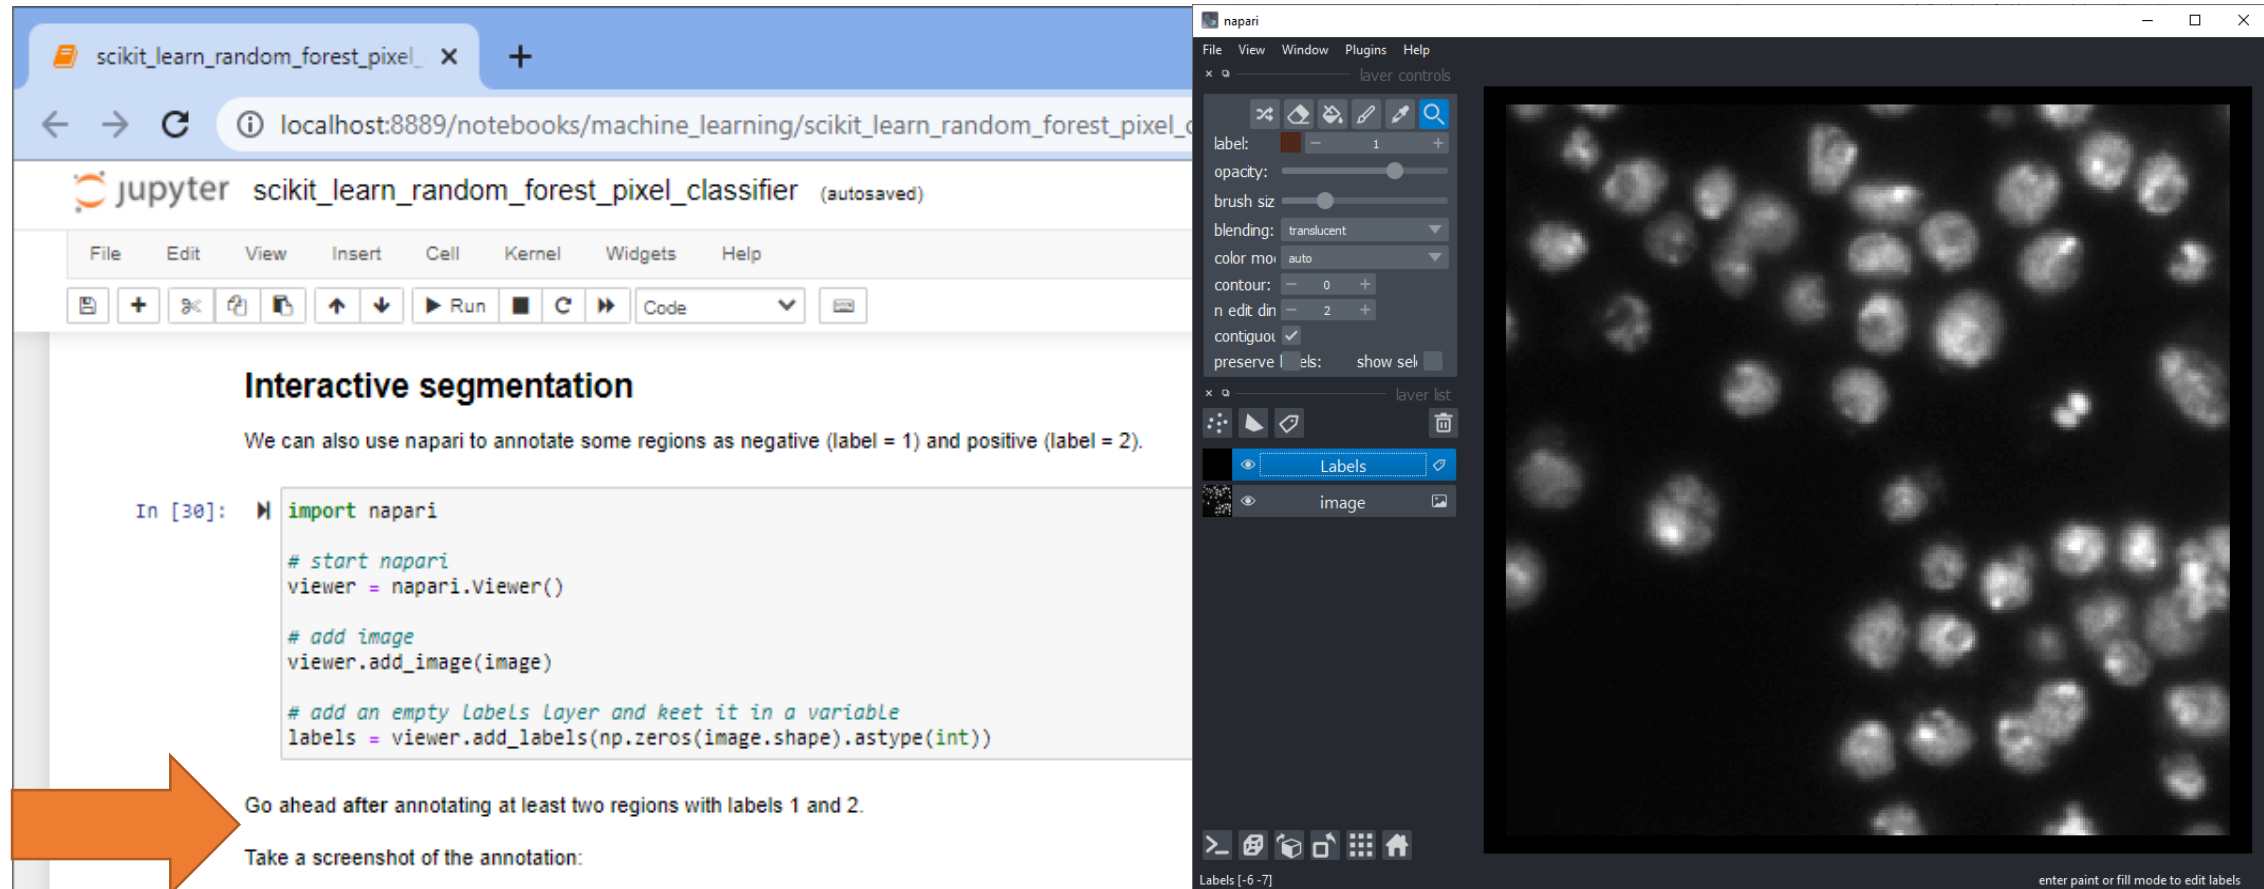


Image data source: [BBBC038v1](https://github.com/BiAPoL/Bio-image_Analysis_with_Python/blob/main/09_machine_learning/02_scikit_learn_random_forest_pixel_classifier.ipynb), available from the Broad Bioimage Benchmark Collection (Caicedo et al., Nature Methods, 2019).  
[https://github.com/BiAPoL/Bio-image\\_Analysis\\_with\\_Python/blob/main/09\\_machine\\_learning/02\\_scikit\\_learn\\_random\\_forest\\_pixel\\_classifier.ipynb](https://github.com/BiAPoL/Bio-image_Analysis_with_Python/blob/main/09_machine_learning/02_scikit_learn_random_forest_pixel_classifier.ipynb)

# Interactive pixel classification

- Jupyter notebooks and napari side-by-side



The image shows a side-by-side view of a Jupyter notebook and the napari viewer. The Jupyter notebook on the left is titled 'scikit\_learn\_random\_forest\_pixel\_classifier' and contains a code cell with the following Python code:

```
In [30]: import napari

# start napari
viewer = napari.Viewer()

# add image
viewer.add_image(image)

# add an empty Labels layer and keep it in a variable
labels = viewer.add_labels(np.zeros(image.shape).astype(int))
```

Below the code cell, an orange arrow points to the text: "Go ahead after annotating at least two regions with labels 1 and 2. Take a screenshot of the annotation:". The napari viewer on the right displays a grayscale image of cell nuclei. The 'layer controls' panel on the left of the viewer shows the 'Labels' layer selected, with settings for label (1), opacity, brush size, blending (translucent), color mode (auto), contour (0), n edit dim (2), and contiguous checked. The 'layer list' on the right shows the 'Labels' layer and the 'image' layer. The status bar at the bottom of the viewer indicates 'Labels [-6 -7]' and 'enter paint or fill mode to edit labels'.

Image data source: [BBC038v1](https://www.ebi.ac.uk/biocompare/benchmark/BBC038v1), available from the Broad Bioimage Benchmark Collection (Caicedo et al., Nature Methods, 2019).

[https://github.com/BiAPoL/Bio-image\\_Analysis\\_with\\_Python/blob/main/09\\_machine\\_learning/02\\_scikit\\_learn\\_random\\_forest\\_pixel\\_classifier.ipynb](https://github.com/BiAPoL/Bio-image_Analysis_with_Python/blob/main/09_machine_learning/02_scikit_learn_random_forest_pixel_classifier.ipynb)

# Interactive pixel classification

- Jupyter notebooks and napari side-by-side

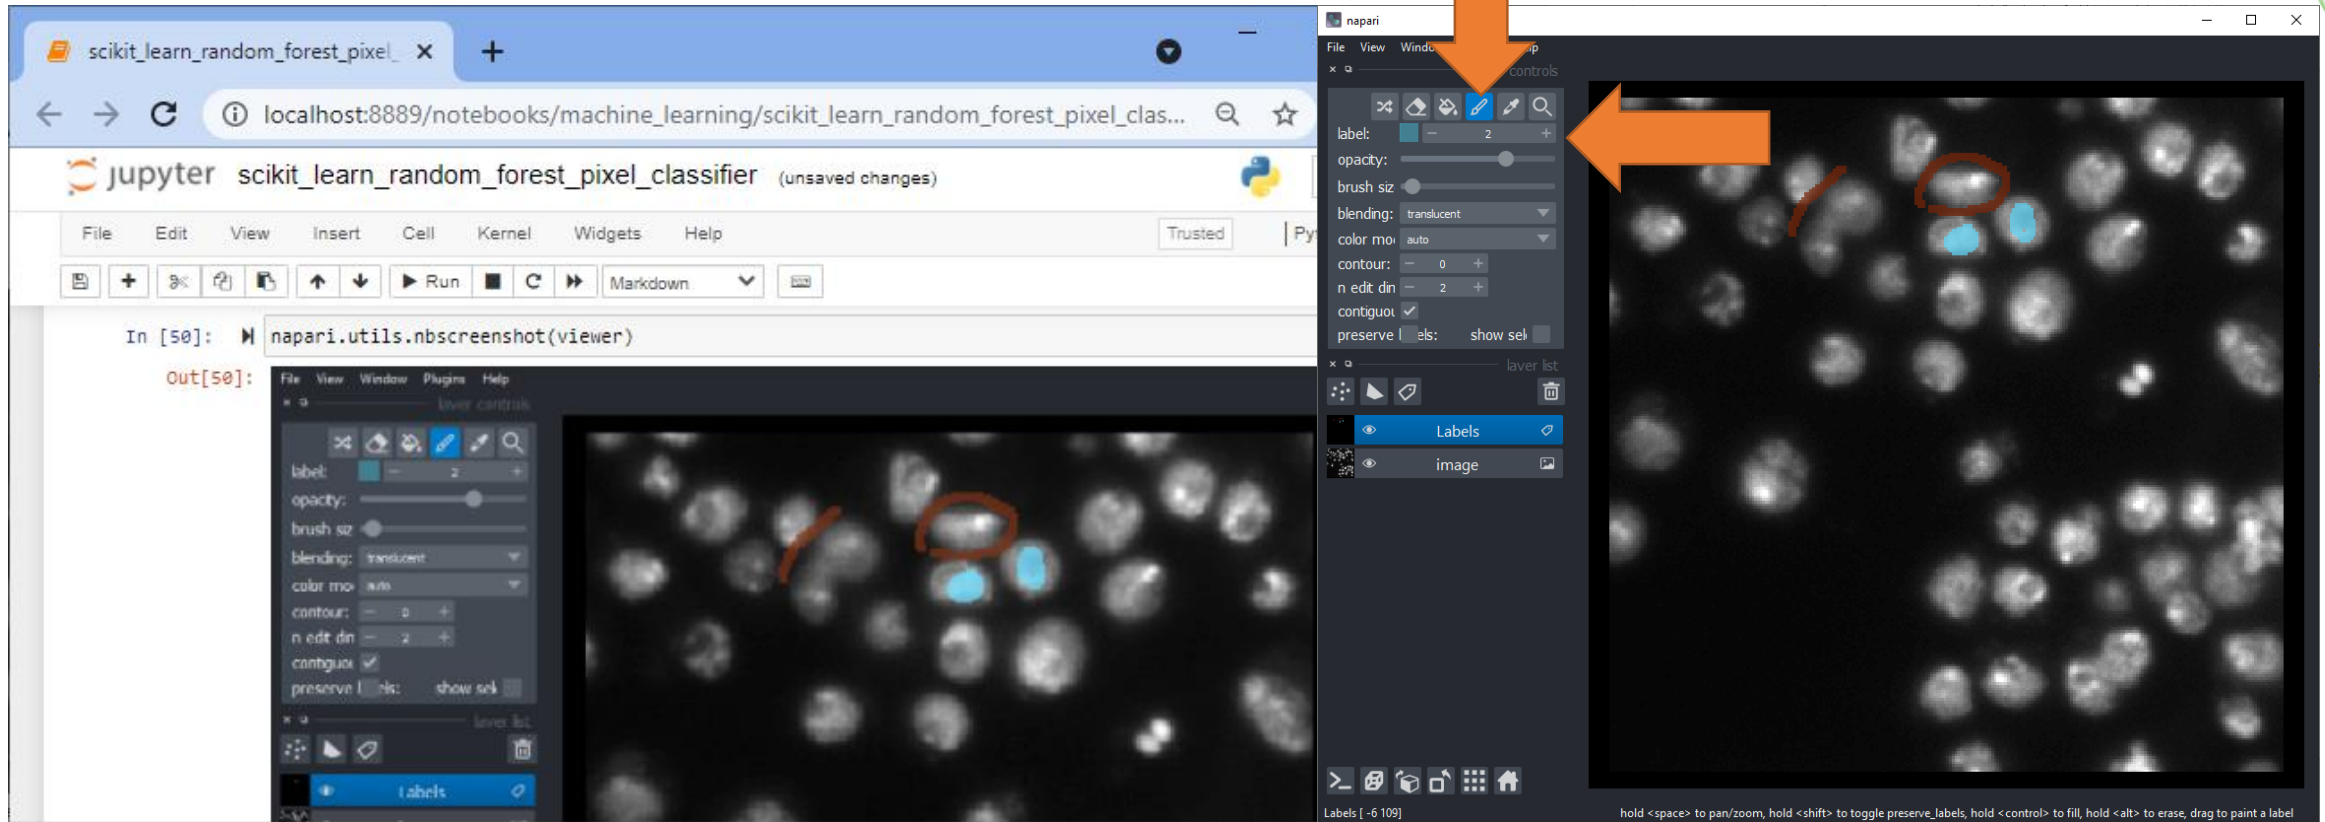


Image data source: [BBBC038v1](https://github.com/BiAPoL/Bio-image_Analysis_with_Python/blob/main/09_machine_learning/02_scikit_learn_random_forest_pixel_classifier.ipynb), available from the Broad Bioimage Benchmark Collection (Caicedo et al., Nature Methods, 2019).

[https://github.com/BiAPoL/Bio-image\\_Analysis\\_with\\_Python/blob/main/09\\_machine\\_learning/02\\_scikit\\_learn\\_random\\_forest\\_pixel\\_classifier.ipynb](https://github.com/BiAPoL/Bio-image_Analysis_with_Python/blob/main/09_machine_learning/02_scikit_learn_random_forest_pixel_classifier.ipynb)



# Interactive pixel classification

- Jupyter notebooks and napari side-by-side

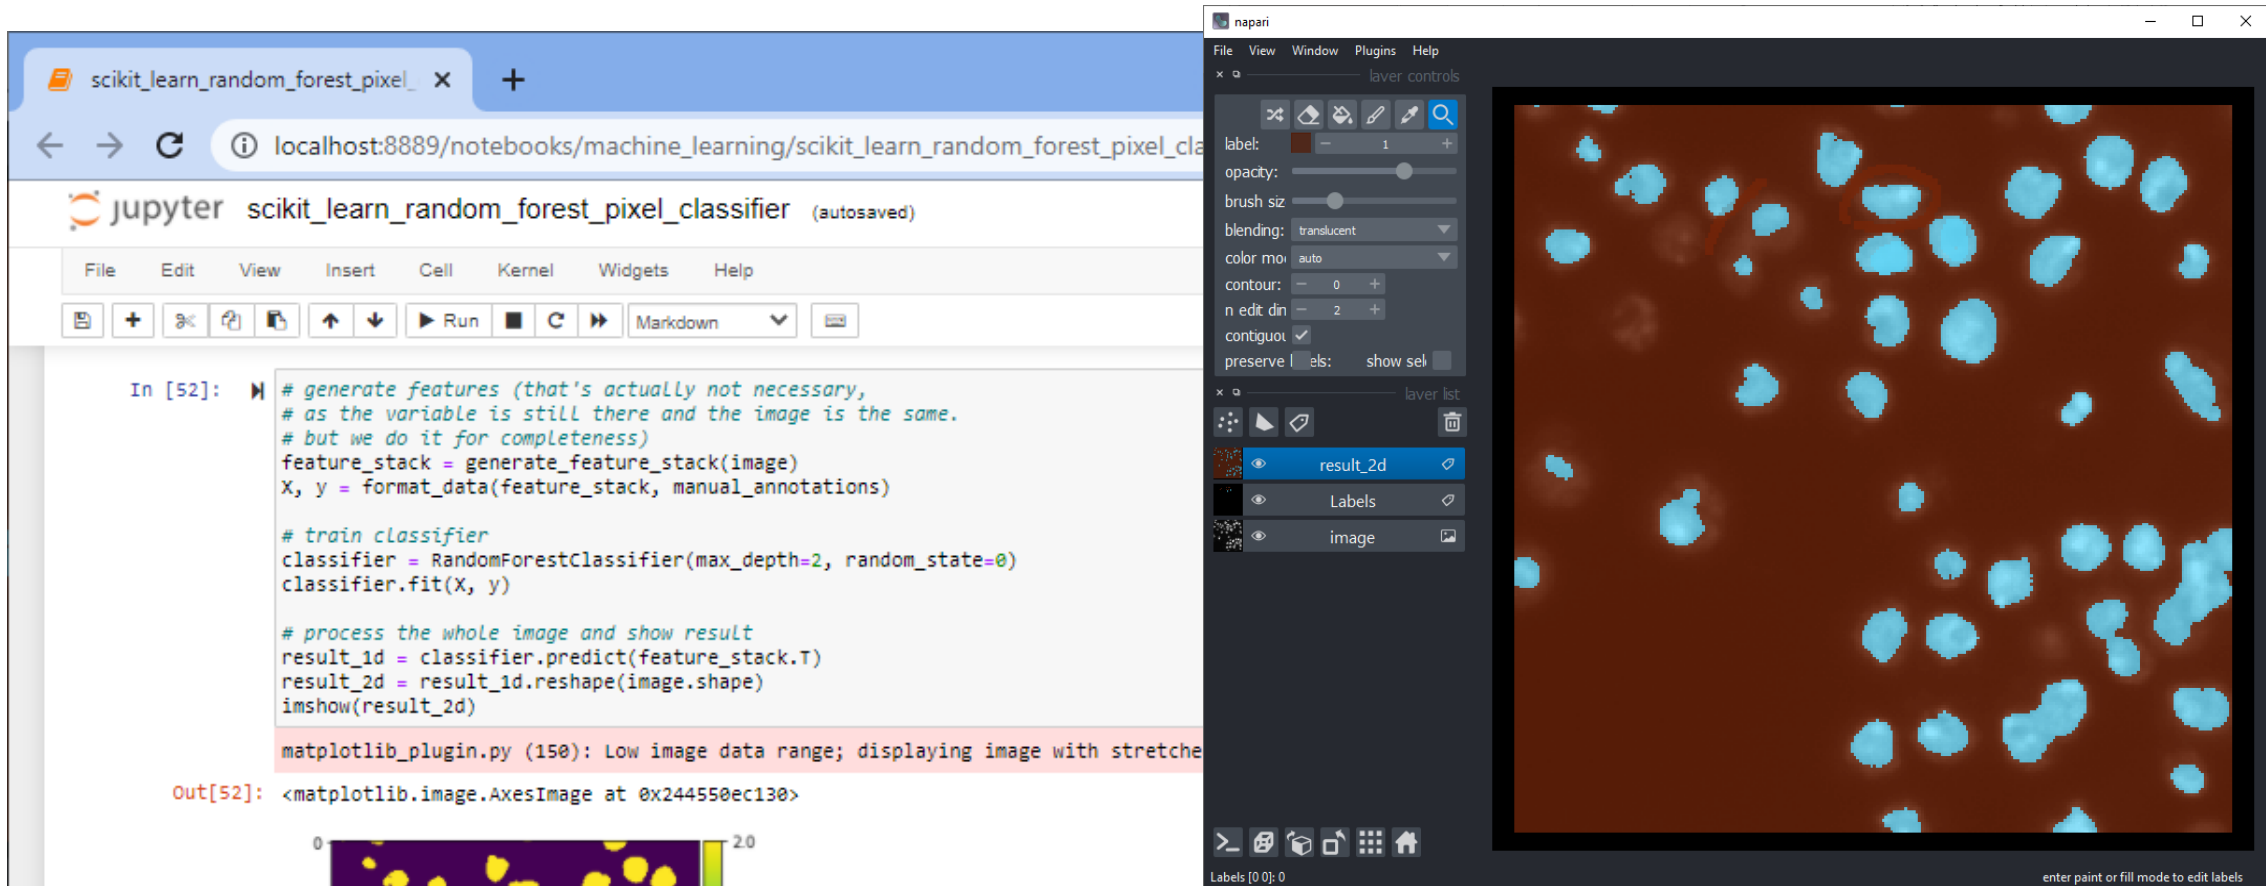


Image data source: [BBBC038v1](https://bbbc038v1), available from the Broad Bioimage Benchmark Collection (Caicedo et al., Nature Methods, 2019).

[https://github.com/BiAPoL/Bio-image\\_Analysis\\_with\\_Python/blob/main/09\\_machine\\_learning/02\\_scikit\\_learn\\_random\\_forest\\_pixel\\_classifier.ipynb](https://github.com/BiAPoL/Bio-image_Analysis_with_Python/blob/main/09_machine_learning/02_scikit_learn_random_forest_pixel_classifier.ipynb)



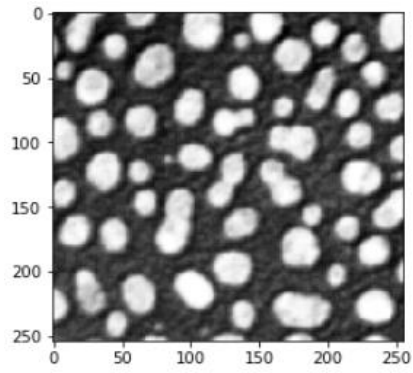


# Accelerated pixel and object classification (APOC)

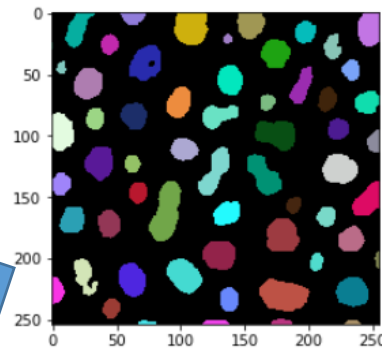
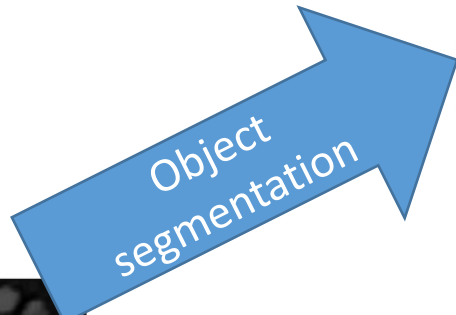
With material from  
Robert Haase

# Accelerated pixel and object classification

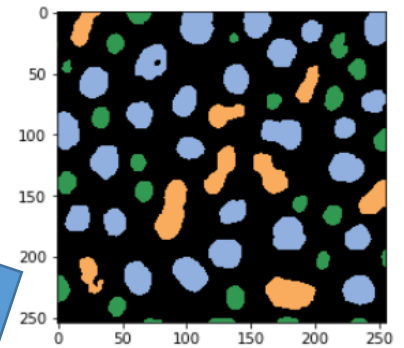
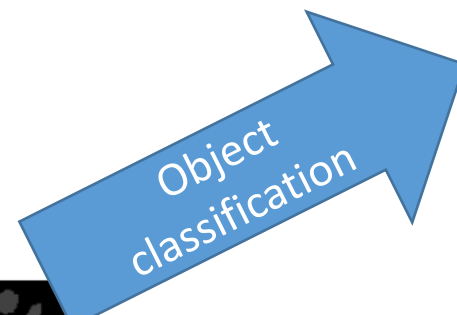
- APOC is a python library that makes use of OpenCL-compatible Graphics Cards to accelerate pixel and object classification



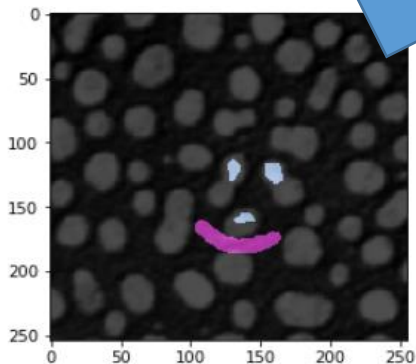
Raw image



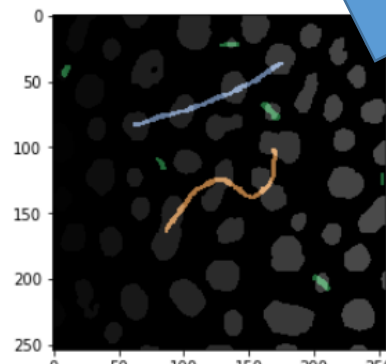
Object label image



Class label image



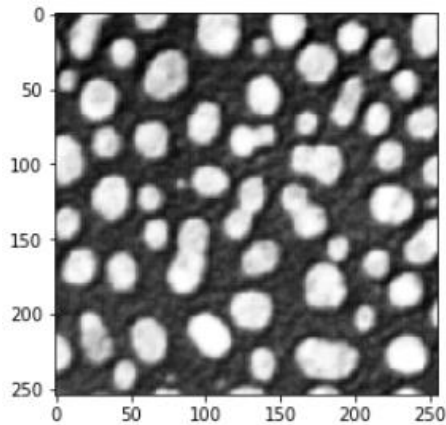
Pixel annotation



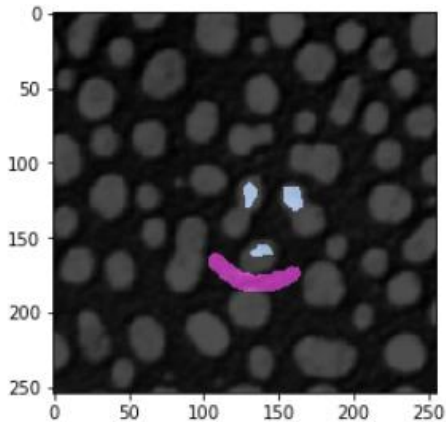
Object annotation

# Object segmentation

- Pixel classification + connected component labeling



Raw image



Pixel annotation

```
# define features
features = "gaussian_blur=1 gaussian_blur=5 sobel_of_gaussian_blur=1"

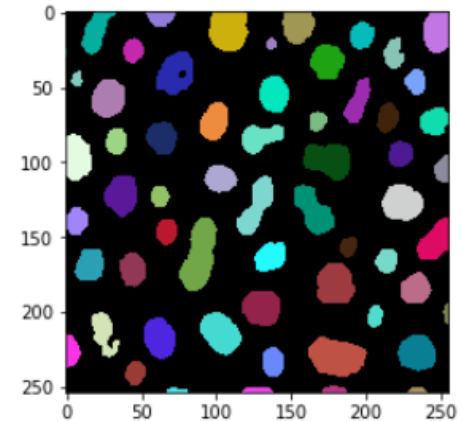
# this is where the model will be saved
cl_filename = 'my_object_segmenter.cl'

# delete classifier in case the file exists already
apoc.erase_classifier(cl_filename)

# train classifier
clf = apoc.ObjectSegmenter(opencv_filename=cl_filename, positive_class_identifier=2)
clf.train(features, manual_annotations, image)

segmentation_result = clf.predict(features=features, image=image)
cle.imshow(segmentation_result, labels=True)
```

Object segmentation

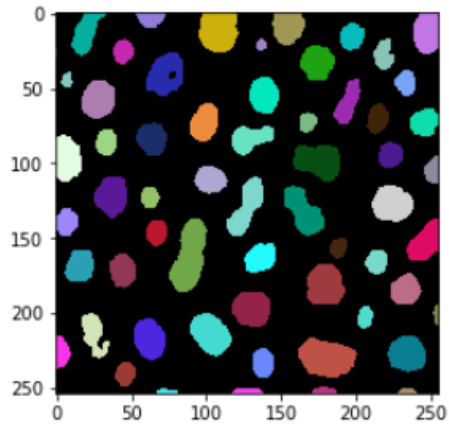


Object label image

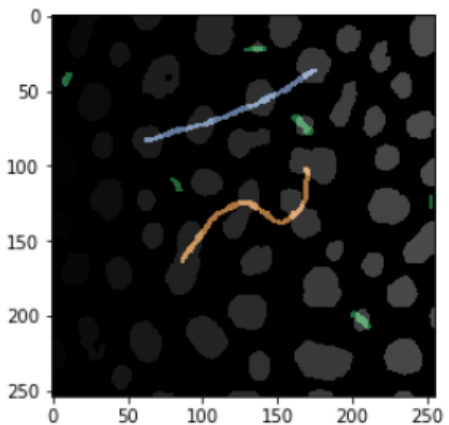
[https://github.com/BiAPoL/Bio-image\\_Analysis\\_with\\_Python/blob/main/09\\_machine\\_learning/03\\_apoc\\_object\\_segmenter.ipynb](https://github.com/BiAPoL/Bio-image_Analysis_with_Python/blob/main/09_machine_learning/03_apoc_object_segmenter.ipynb)

# Object classification

- Feature extraction + tabular classification



Object label image



Object annotation

```
# for the classification we define size and shape as criteria  
features = 'area mean_max_distance_to_centroid_ratio'
```

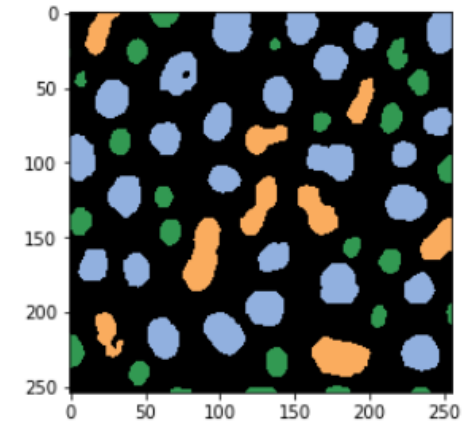
```
# This is where the model will be saved  
cl_filename_object_classifier = "my_object_classifier.cl"
```

```
# delete classifier in case the file exists already  
apoc.erase_classifier(cl_filename_object_classifier)
```

```
# train the classifier  
classifier = apoc.ObjectClassifier(cl_filename_object_classifier)  
classifier.train(features, segmentation_result, annotation, image)
```

```
# determine object classification  
classification_result = classifier.predict(segmentation_result, image)  
cle.imshow(classification_result, labels=True)
```

Object classification

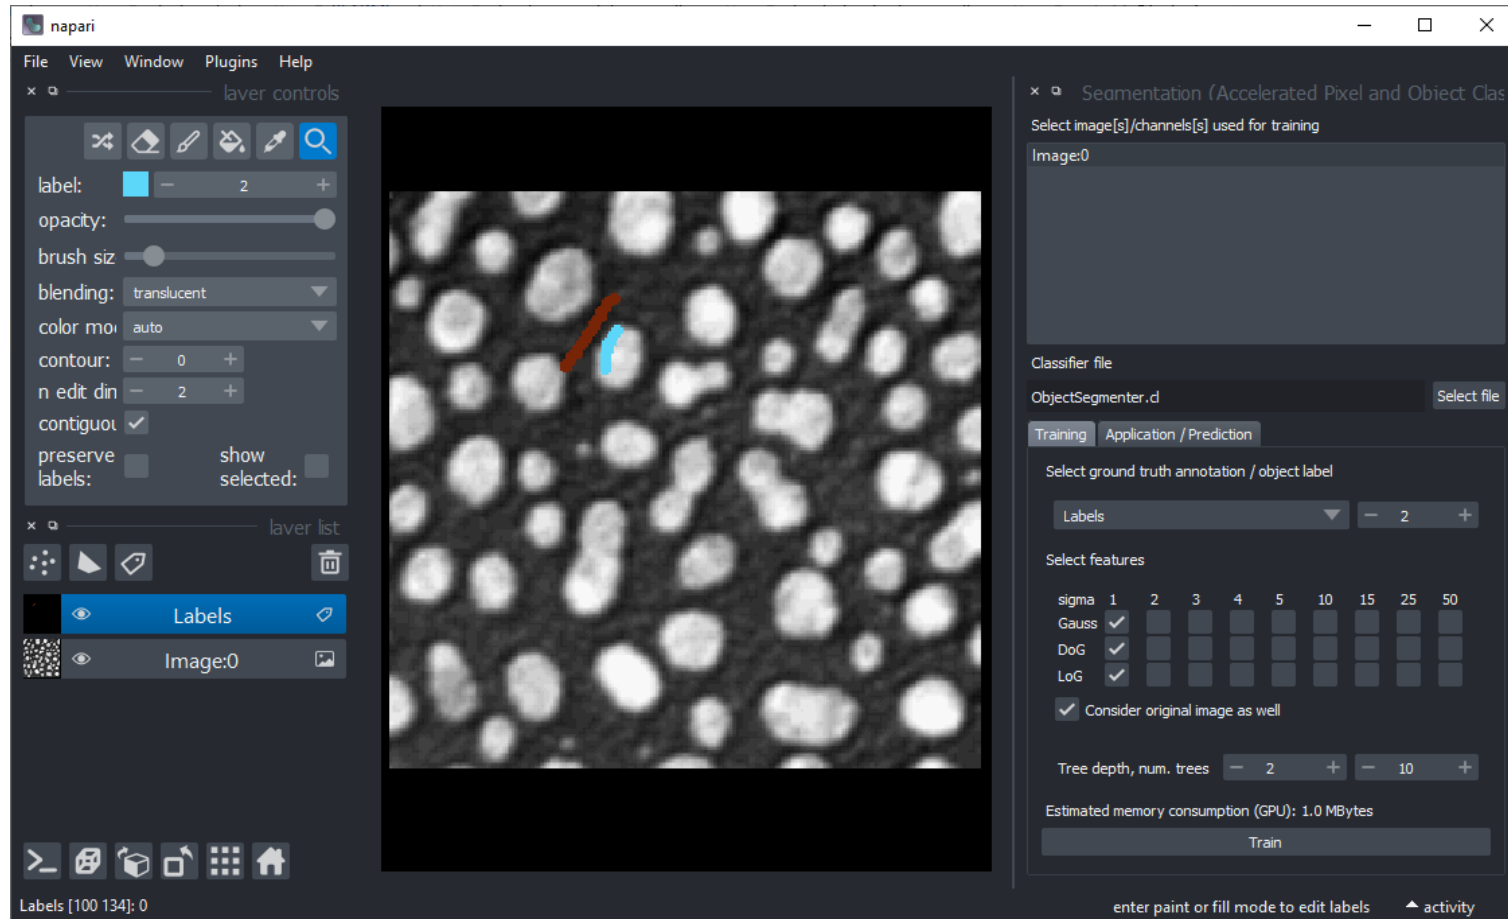


Class label image

[https://github.com/BiAPoL/Bio-image\\_Analysis\\_with\\_Python/blob/main/09\\_machine\\_learning/03\\_apoc\\_object\\_segementer.ipynb](https://github.com/BiAPoL/Bio-image_Analysis_with_Python/blob/main/09_machine_learning/03_apoc_object_segementer.ipynb)

# Graphical user interface

- Object segmentation
- <https://github.com/haesleinhuepf/napari-accelerated-pixel-and-object-classification#object-and-semantic-segmentation>





# Supervised machine learning for tissue classification

Random Forest  
Classifiers based on

- scikit-learn and
- clesperanto

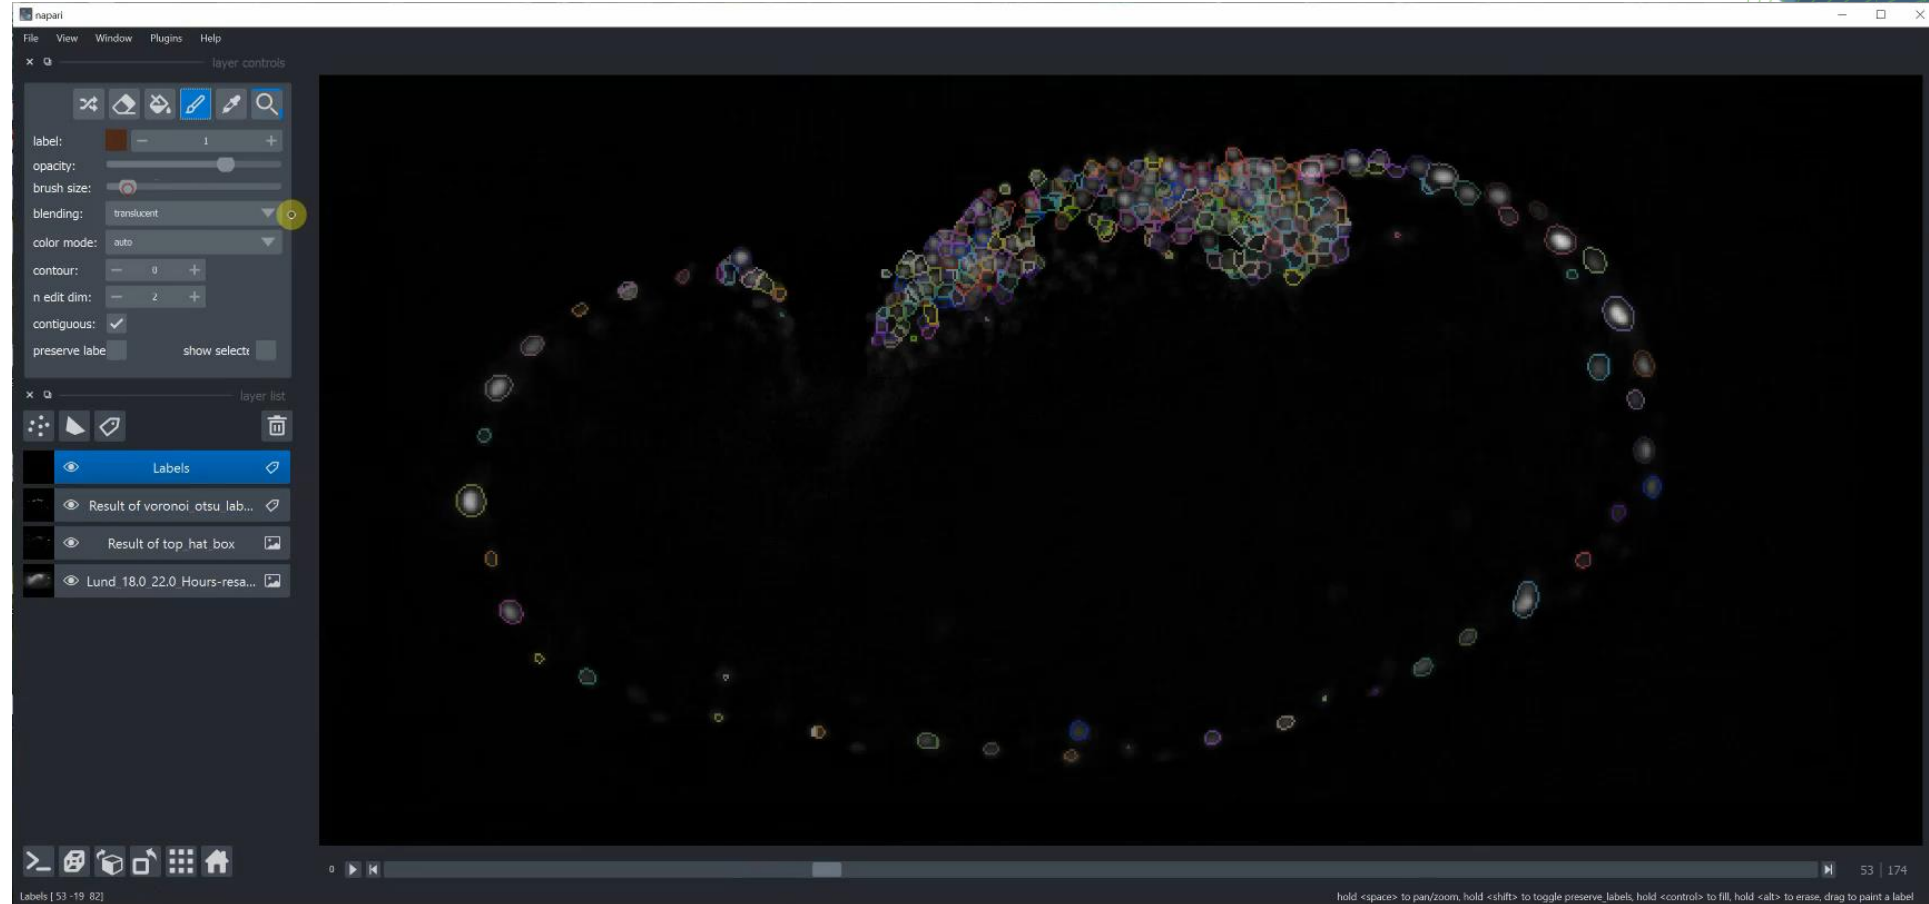


Image data source: Daniela Vorkel, Myers lab, MPI-CBG/CSBD

<https://github.com/haesleinhuepf/napari-accelerated-pixel-and-object-classification>



# Data exploration / supervised machine learning

- Inspect how the random forest classifier makes decisions
- Note: Beware of correlated parameters!

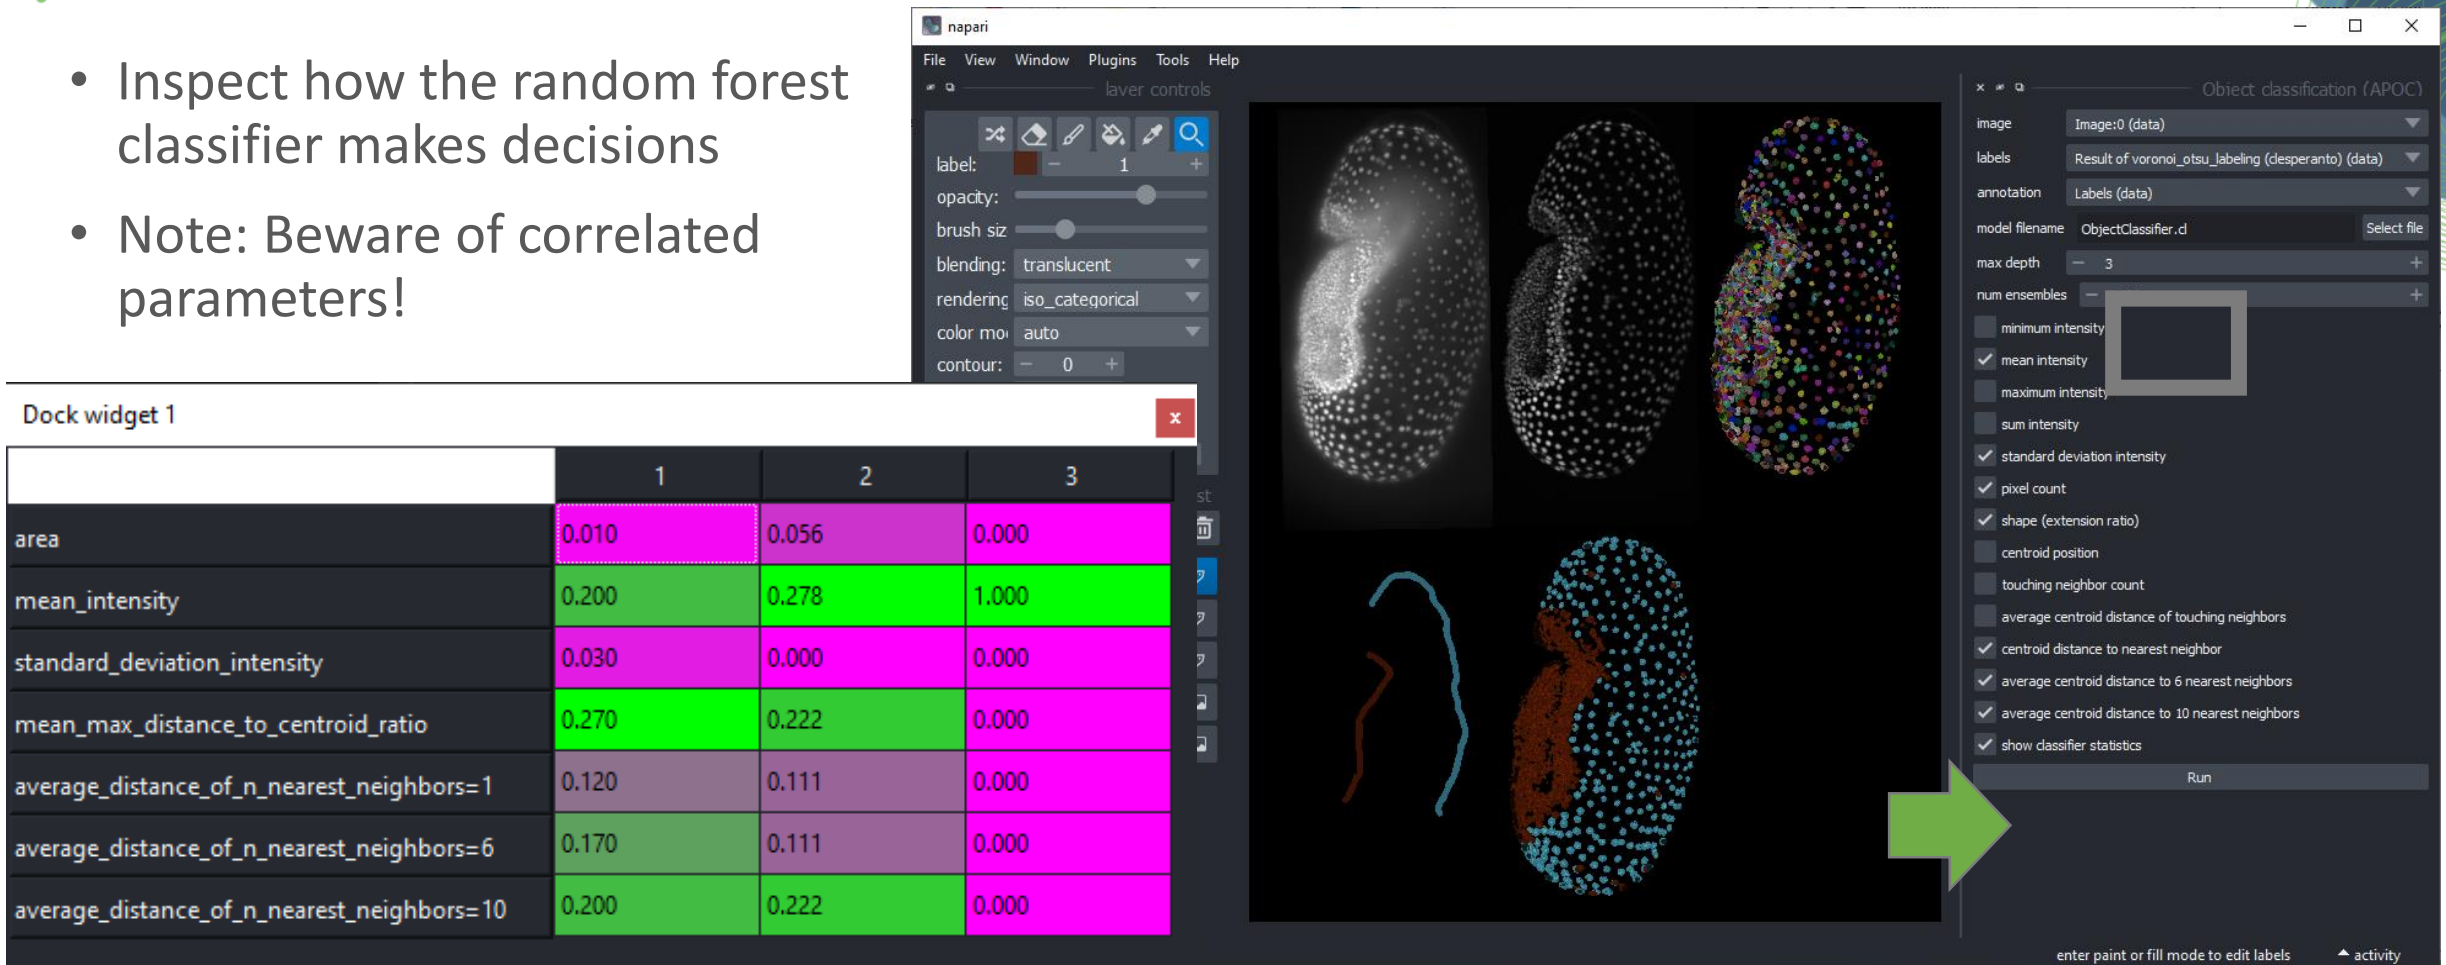


Image data source: Daniela Vorkel, Myers lab, MPI-CBG/CSBD  
<https://github.com/haesleinhuepf/napari-accelerated-pixel-and-object-classification>

# Data exploration / supervised machine learning

- Inspect how the random forest classifier makes decisions
- Note: Beware of correlated parameters!

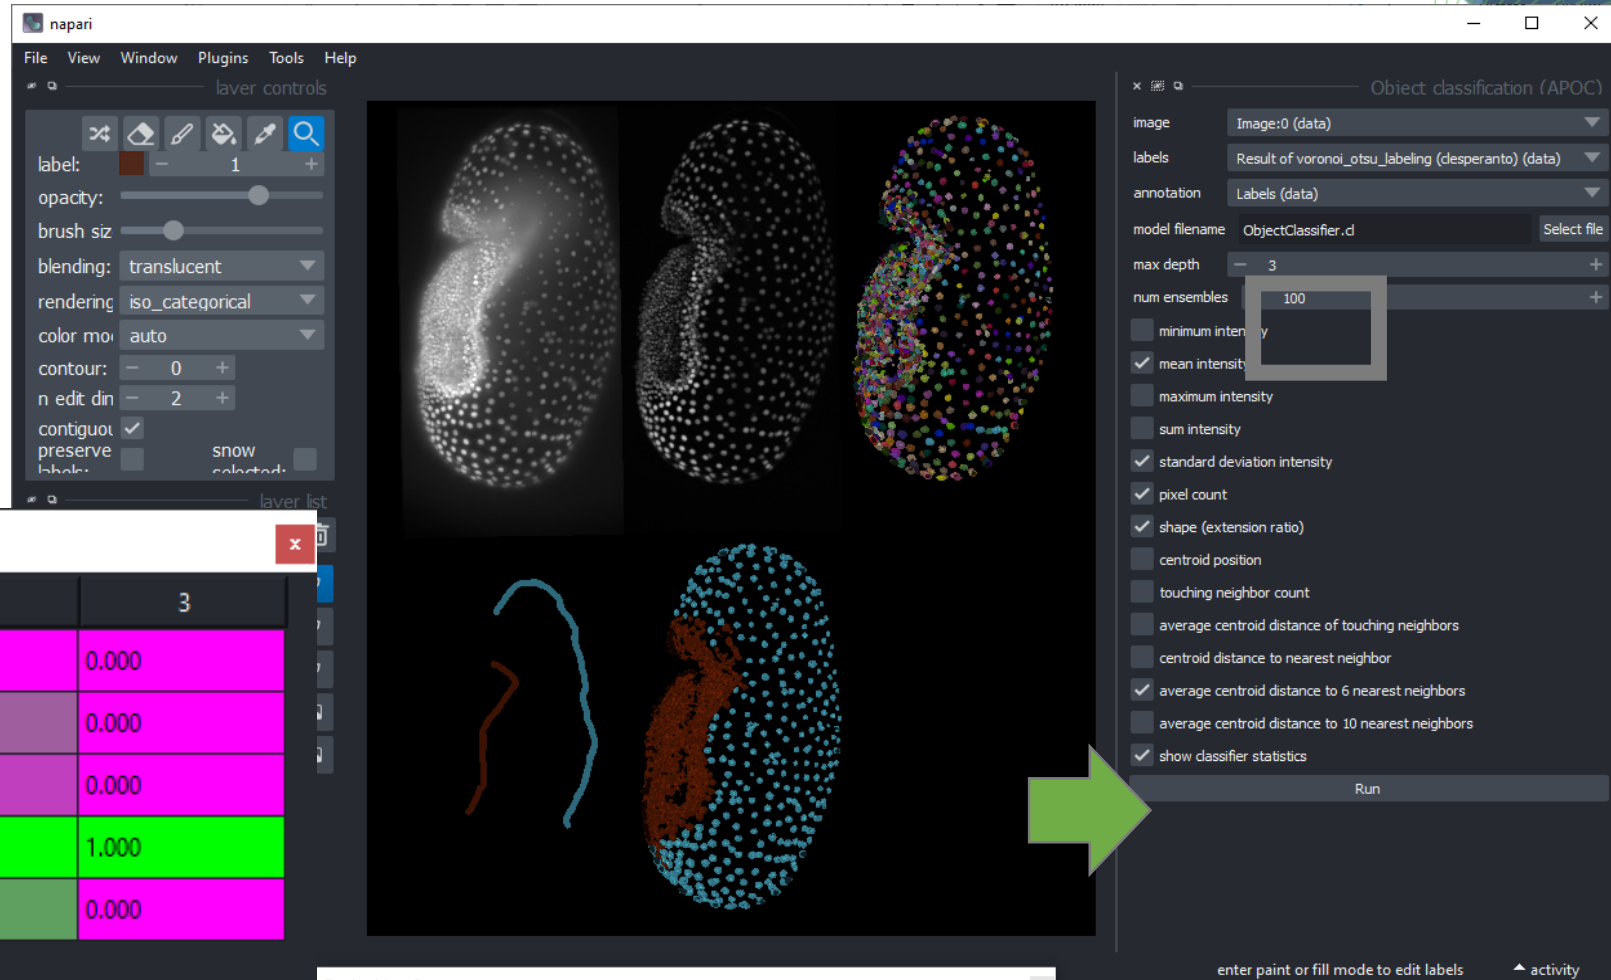


Image data source: Daniela Vorkel, Myers lab, MPI-CBG/CSBD

<https://github.com/haesleinhuepf/napari-accelerated-pixel-and-object-classification>

# Data exploration / supervised machine learning

- Inspect how the random forest classifier makes decisions
- Note: Beware of correlated parameters!

Dock widget 2

	1	2	3
area	0.060	0.000	0.000
mean_intensity	0.330	0.167	0.000
standard_deviation_intensity	0.040	0.111	0.000
mean_max_distance_to_centroid_ratio	0.260	0.444	1.000
average_distance_of_n_nearest_neighbors=6	0.310	0.278	0.000

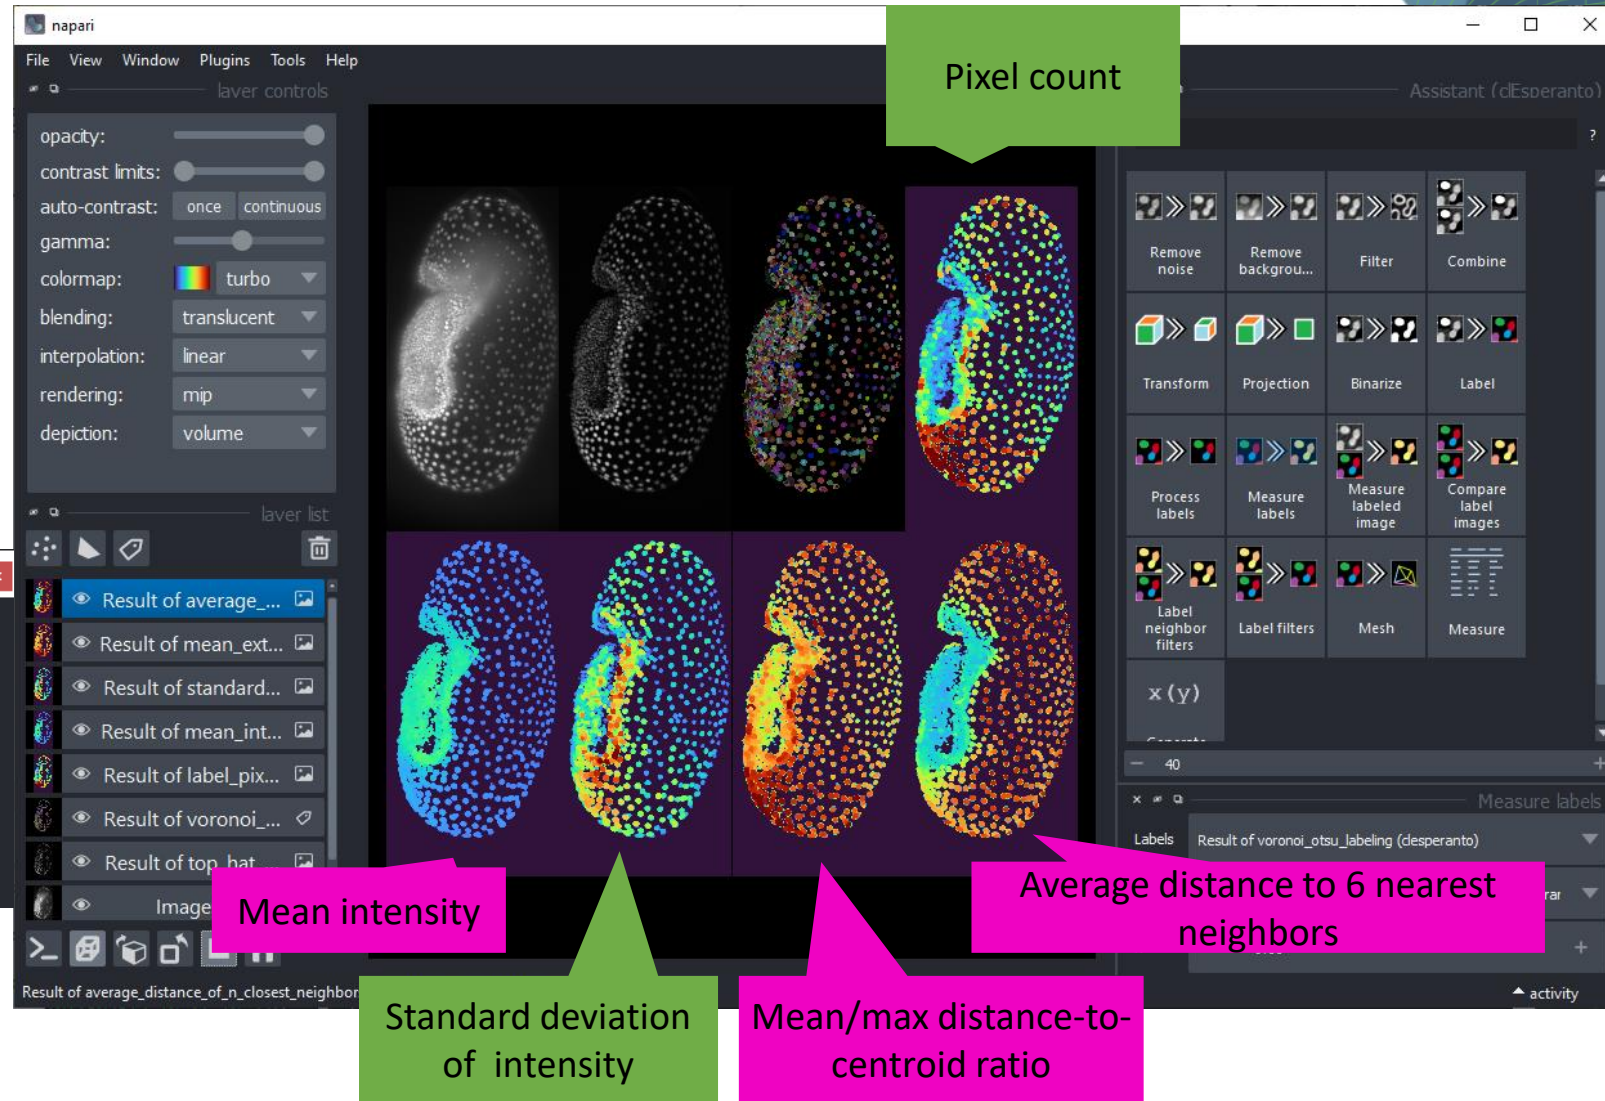
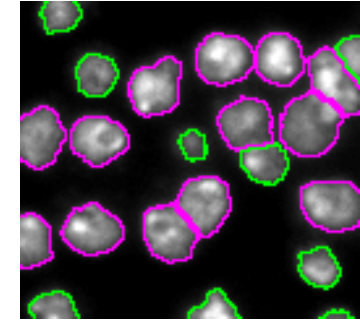
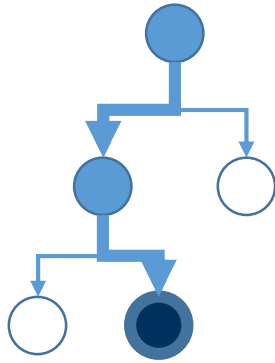


Image data source: Daniela Vorkel, Myers lab, MPI-CBG/CSBD  
[https://github.com/clEsperanto/napari\\_pyclesperanto\\_assistant](https://github.com/clEsperanto/napari_pyclesperanto_assistant)



# Thank you for your attention!



with material from

Robert Haase, ScaDS.AI, Leipzig University

Deborah Schmidt, Jug Lab, MPI CBG

Uwe Schmidt, Myers Lab, MPI CBG

Martin Weigert, EPFL

Ignacio Arganda-Carreras, Universidad del Pais Vasco

Carsen Stringer, HHMI Janelia

Wei Ouyang, KTH Royal Institute of Technology, Stockholm and

The Scikit-Learn community