

Bio-Image Data Science Training

TRAINING: Tabular Data Wrangling

SPEAKER: Matthias Täschner

Including material from Robert Haase

GEFÖRDERT VOM



Bundesministerium
für Bildung
und Forschung



SACHSEN Diese Maßnahme wird gefördert durch die Bundesregierung aufgrund eines Beschlusses des Deutschen Bundestages. Diese Maßnahme wird mitfinanziert durch Steuermittel auf der Grundlage des von den Abgeordneten des Sächsischen Landtags beschlossenen Haushaltes.





AGENDA

- pandas for tabular data
- DataFrame and Series
- Creating DataFrames
- Selecting from DataFrames
 - Label-location
 - Integer-location
 - Boolean indexing
- Combining DataFrames
- Handle Missing Data
- Tidy Data

pandas for tabular data

[pandas](#) is an open-source Python library for data manipulation and analysis

- Development started in 2008, recent stable version is 2.2.2
- Offers data structures and operations tailored for tabular data and time series analysis
- Core structures
 - [Series](#) (1-dimensional labeled array with index, i.e., a one-column table)
 - [DataFrame](#) (2-dimensional data structure with labels and index)
- Built on top of
 - [NumPy](#) for performant numerical operations in Python
 - [SciPy](#) for scientific computations in Python
 - [Matplotlib](#) for data visualization in Python
- A good starting point is the [10 minutes to pandas](#) tutorial



© pandas via NumFOCUS, Inc.
<https://pandas.pydata.org/>

DataFrame and Series

Index with labels for columns

		0	1	2	3	4	5
		Attr1	Attr2	Attr3	Attr4	TimeStamp	Valid
0	A	filename12	15.5	188	0.02	2024-05-15	True
1	B	filename45	36.5	211	0.1	2024-04-12	True
2	C	filename98	80.1	119	0.65	2023-12-11	False
3	D	filename32	24.7	75	0.08	2024-02-25	True

Row axis

Data

Column axis

Index with labels for rows

Series

Uses the rows index

Creating DataFrames

- DataFrames can be created from different inputs
- Let's assume we have the following data available from measurements

sample	area	minor_axis	major_axis
A	45	2	3
B	23	4	4
C	68	4	4

```
1 # Import the pandas module
2 import pandas as pd
Executed at 2024.05.07 08:11:23 in 254ms
```

```
1 # Measurements as dict of lists
2 measurements = {
3     "sample": ['A', 'B', 'C'],
4     "area": [45, 23, 68],
5     "minor_axis": [2, 4, 4],
6     "major_axis": [3, 4, 5],
7 }
8 # Create DataFrame from dict
9 df = pd.DataFrame(data=measurements)
10 print(df)
Executed at 2024.05.07 08:16:49 in 5ms
```

```
sample  area  minor_axis  major_axis
0      A    45           2           3
1      B    23           4           4
2      C    68           4           5
```

```
1 # Measurements as nested lists
2 data = [
3     ['A', 'B', 'C'],
4     [45, 23, 68],
5     [2, 4, 4],
6     [3, 4, 5],
7 ]
8 # We need to provide the labels here
9 labels = ['sample', 'area', 'minor_axis', 'major_axis']
10 # Create DataFrame from nested lists with separate labels
11 df = pd.DataFrame(data=data, index=labels)
12 print(df)
Executed at 2024.05.07 08:25:31 in 5ms
```

```
sample  A  B  C
area    45 23 68
minor_axis  2 4 4
major_axis  3 4 5
```

```
1 # Oops, it's rotated.
2 # We can fix this via transposing
3 df = df.transpose()
4 print(df)
Executed at 2024.05.07 08:37:30 in 4ms
```

```
sample area minor_axis major_axis
0      A    45           2           3
1      B    23           4           4
2      C    68           4           5
```

This has another format??

Creating DataFrames

- DataFrames provide some convenient methods to get an overview on the structure and data

```
df.info()
```

Executed at 2024.05.07 08:51:10 in 8ms

Info about the
columns data type

Info about the
rows index

Info about the
columns

Info about the
columns index

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3 entries, 0 to 2
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   sample      3 non-null      string
1   area        3 non-null      Int64
2   minor_axis  3 non-null      Int64
3   major_axis  3 non-null      Int64
dtypes: Int64(3), string(1)
memory usage: 233.0 bytes
```

First info about
missing data

```
df.describe(include='all')
```

Executed at 2024.05.07 08:45:36 in 9ms

Show descriptive
statistics for...

	sample	area	minor_axis	major_axis
count	3	3.0	3.0	3.0
unique	3	<NA>	<NA>	<NA>
top	A	<NA>	<NA>	<NA>
freq	1	<NA>	<NA>	<NA>
mean	NaN	45.333333	3.333333	4.0
std	NaN	22.501852	1.154701	1.0
min	NaN	23.0	2.0	3.0
25%	NaN	34.0	3.0	3.5
50%	NaN	45.0	4.0	4.0
75%	NaN	56.5	4.0	4.5
max	NaN	68.0	4.0	5.0

...categorical
data

...numerical
data

Selecting from DataFrames

Overview

There are different ways to select data from a DataFrame

Operation	Syntax	Result
Select one column	<code>df[column label]</code>	Series
Select several columns	<code>df[list of column labels]</code>	DataFrame
Select one row by label	<code>df.loc[row label]</code>	Series
Select several rows by label	<code>df.loc[list of row labels]</code>	DataFrame
Select rows by slicing on labels	<code>df.loc[start row label : end row label]</code>	DataFrame
Select one row by integer location	<code>df.iloc[index number]</code>	Series
Select several rows by integer location	<code>df.iloc[list of index numbers]</code>	DataFrame
Select rows by slicing on integer location	<code>df.iloc[start index number : end index number]</code>	DataFrame
Select rows and columns by label	<code>df.loc[list of row labels , list of column labels]</code>	DataFrame
... there are more possibilities		

Selecting from DataFrames

Label-location

Here, we use the columns' labels and rows' labels to select data – `[]` and `loc[]`

```
1 # Create DataFrame from dict
2 df = pd.DataFrame(data=measurements)
3 print(df)
```

Executed at 2024.05.07 10:18:39 in 4ms

	sample	area	minor_axis	major_axis
0	A	45	2	3
1	B	23	4	4
2	C	68	4	5

```
1 # Select data for column 'area'
2 print(df['area'])
```

Executed at 2024.05.07 10:19:42 in 2ms

0	45
1	23
2	68

Name: area, dtype: int64

```
1 # Select data for columns 'area' and 'minor_axis'
2 print(df[['area', 'minor_axis']])
```

Executed at 2024.05.07 10:20:44 in 4ms

	area	minor_axis
0	45	2
1	23	4
2	68	4

```
1 # Select data for row with label 2
2 print(df.loc[2])
```

Executed at 2024.05.07 10:23:01 in 3ms

sample	C
area	68
minor_axis	4
major_axis	5

Name: 2, dtype: object

```
1 # Select data for rows with label 2 and 0
2 print(df.loc[[2, 0]])
```

Executed at 2024.05.07 10:24:44 in 5ms

	sample	area	minor_axis	major_axis
2	C	68	4	5
0	A	45	2	3

Selecting from DataFrames

Integer-location

Here, we use the integer indices of the rows and columns to select data – [iloc\[\]](#)

```
1 # Create DataFrame from dict
2 df = pd.DataFrame(data=measurements)
3 print(df)
```

Executed at 2024.05.07 10:18:39 in 4ms

	sample	area	minor_axis	major_axis
0	A	45	2	3
1	B	23	4	4
2	C	68	4	5

```
1 # Select the first row (index at 0)
2 print(df.iloc[0])
```

Executed at 2024.05.07 15:39:49 in 3ms

	sample	area	minor_axis	major_axis
0	A	45	2	3

Name: 0, dtype: object

```
1 # Select all rows starting at index 1
2 print(df.iloc[1:])
```

Executed at 2024.05.07 15:41:24 in 5ms

	sample	area	minor_axis	major_axis
1	B	23	4	4
2	C	68	4	5

```
1 # Select rows with index 0 and 2
2 # ...and the column with index 0
3 print(df.iloc[[0,2], 0])
```

Executed at 2024.05.07 15:44:36 in 4ms

	sample
0	A
2	C

Name: sample, dtype: string

```
1 # Select all rows (via :)
2 # ...and all columns starting from index 2
3 print(df.iloc[:, 2:])
```

Executed at 2024.05.07 15:46:19 in 3ms

	minor_axis	major_axis
0	2	3
1	4	4
2	4	5

Selecting from DataFrames

Boolean indexing

Here, we use conditional / logical queries (masking) to select data – [Boolean vectors](#)

```
1 # Create DataFrame from dict
2 df = pd.DataFrame(data=measurements)
3 print(df)
```

Executed at 2024.05.07 10:18:39 in 4ms

	sample	area	minor_axis	major_axis
0	A	45	2	3
1	B	23	4	4
2	C	68	4	5

```
1 # Select all data where 'area' is greater than 50
2 criterion = df['area'] > 50
3 print(df[criterion])
```

Executed at 2024.05.07 15:54:59 in 5ms

	sample	area	minor_axis	major_axis
2	C	68	4	5

```
1 # Select all data where minor_axis is not 4
2 print(df[df['minor_axis'] != 4])
```

Executed at 2024.05.07 16:02:37 in 3ms

	sample	area	minor_axis	major_axis
0	A	45	2	3

```
1 # Select the sample where 'area' is greater than 50 and minor_axis is 4
2 # Combined expressions must be grouped by using parentheses
3 # We will use loc in combination with boolean vectors
4 print(df.loc[(df['area'] > 50) & (df['minor_axis'] == 4), 'sample'])
```

```
2    C
Name: sample, dtype: string
```

Combining DataFrames

There are 3 main operations to combine DataFrames

- [concat](#) combines an arbitrary number of pandas objects (DataFrames, Series)
 - `pandas.concat([df1, df2])`
- [merge](#) performs SQL-like combination on two pandas objects
 - `pandas.merge(df1, df2)` or `df1.merge(df2)`
- [join](#) is a DataFrame method to combine it with an arbitrary number of other pandas objects
 - `df1.join(df2)` or `df1.join([df2, df3])`

All operations provide parameters for further control, e.g.,

- Along which axis the combination is performed
- What kind of set logic (union or intersection) to use for combination
- On which column (index) a combination is performed
- ...

Handle missing data

Often values are missing, and we need to [detect and handle](#) them

```
1 # Create a DataFrame from dict
2 df = pd.DataFrame(data=measurements)
3 print(df)
```

Executed at 2024.05.07 18:47:12 in 4ms

```
✓
```

	sample	area	minor_axis	major_axis
0	A	45	2	<NA>
1	B	<NA>	4	4
2	C	68	4	<NA>

```
1 # Use info() for a first overview again
2 print(df.info())
```

Executed at 2024.05.07 18:47:45 in 5ms

```
✓
```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3 entries, 0 to 2
Data columns (total 4 columns):

#	Column	Non-Null Count	Dtype
0	sample	3 non-null	string
1	area	2 non-null	Int64
2	minor_axis	3 non-null	Int64
3	major_axis	1 non-null	Int64

pandas' [isnull\(\)](#) provides a Boolean masking for missing values

```
1 # Create masking (True/False) for missing data
2 print(df.isnull())
```

Executed at 2024.05.07 18:52:06 in 4ms

```
✓
```

	sample	area	minor_axis	major_axis
0	False	False	False	True
1	False	True	False	False
2	False	False	False	True

With this masking we can do further checks, and (since True==1 and False==0) even math

```
1 # Check for each column if there are any missing values
2 print(df.isnull().any())
```

Executed at 2024.05.07 19:01:14 in 3ms

```
✓
```

	sample	area	minor_axis	major_axis
	False	True	False	True

dtype: bool

```
1 # Compute and sort percentage of missing values in the columns
2 print(df.isnull().mean().sort_values(ascending=False) * 100)
```

Executed at 2024.05.07 19:03:29 in 4ms

```
✓
```

	major_axis	area	sample	minor_axis
	66.666667	33.333333	0.000000	0.000000

dtype: float64

Handle missing data

But what to do with the knowledge about missing values in the data...

- Ignore and go on with the analysis?
- Remove all samples with missing values from the data?
- Try to fill in the gaps (“data imputation”)?

It depends...
Open for further discussion

Tidy Data

In pandas, you *can* use hierarchical indices (MultiIndex) to “stack” data

Month			January					
Day in month			01			02		
Measurement			Temp	Wind	Pressure	Temp	Wind	Pressure
Country	City	Station						
DE	Leipzig	DE102						
		DE205						
	Berlin	DE035						
		DE962						
GB	London	GB147						
		GB906						
	Bristol	GB781						
		GB006						

Looks nice for humans,
but is tough to analyze

Tidy Data

For data analysis, “tidy data” works better:

- Each variable is a column
- Each observation is a row
- Each type of observation has its own DataFrame

Station	Temp	Wind	Pressure
DE102			
DE205			
DE035			
DE962			
GB147			
GB906			
GB781			
GB006			

Measurements for
January 1st

Date	Temp	Wind	Pressure
2024-01-01			
2024-01-02			
2024-01-03			
2024-01-04			
2024-01-05			
2024-01-06			
2024-01-07			
2024-01-08			

Measurements for
Station DE102

Any questions or remarks?

GEFÖRDERT VOM



Bundesministerium
für Bildung
und Forschung



SACHSEN Diese Maßnahme wird gefördert durch die Bundesregierung aufgrund eines Beschlusses des Deutschen Bundestages. Diese Maßnahme wird mitfinanziert durch Steuermittel auf der Grundlage des von den Abgeordneten des Sächsischen Landtags beschlossenen Haushaltes.

